

Introduction to Stream Processing with Apache Flink

Tugdual Grall
@tgrall



{"about" : "me"}

Tugdual “Tug” Grall

- MapR : Technical Evangelist
 - *MongoDB, Couchbase, eXo, Oracle*
 - *NantesJUG co-founder*
-
- @tgrall
 - <http://tgrall.github.io>
 - tug@mapr.com / tugdual@gmail.com

THE EVOLUTION OF

SOFTWARE ARCHITECTURE

1990's

SPAGHETTI-ORIENTED
ARCHITECTURE
(aka Copy & Paste)



2000's

LASAGNA-ORIENTED
ARCHITECTURE
(aka Layered Monolith)



2010's

RAVIOLI-ORIENTED
ARCHITECTURE
(aka Microservices)

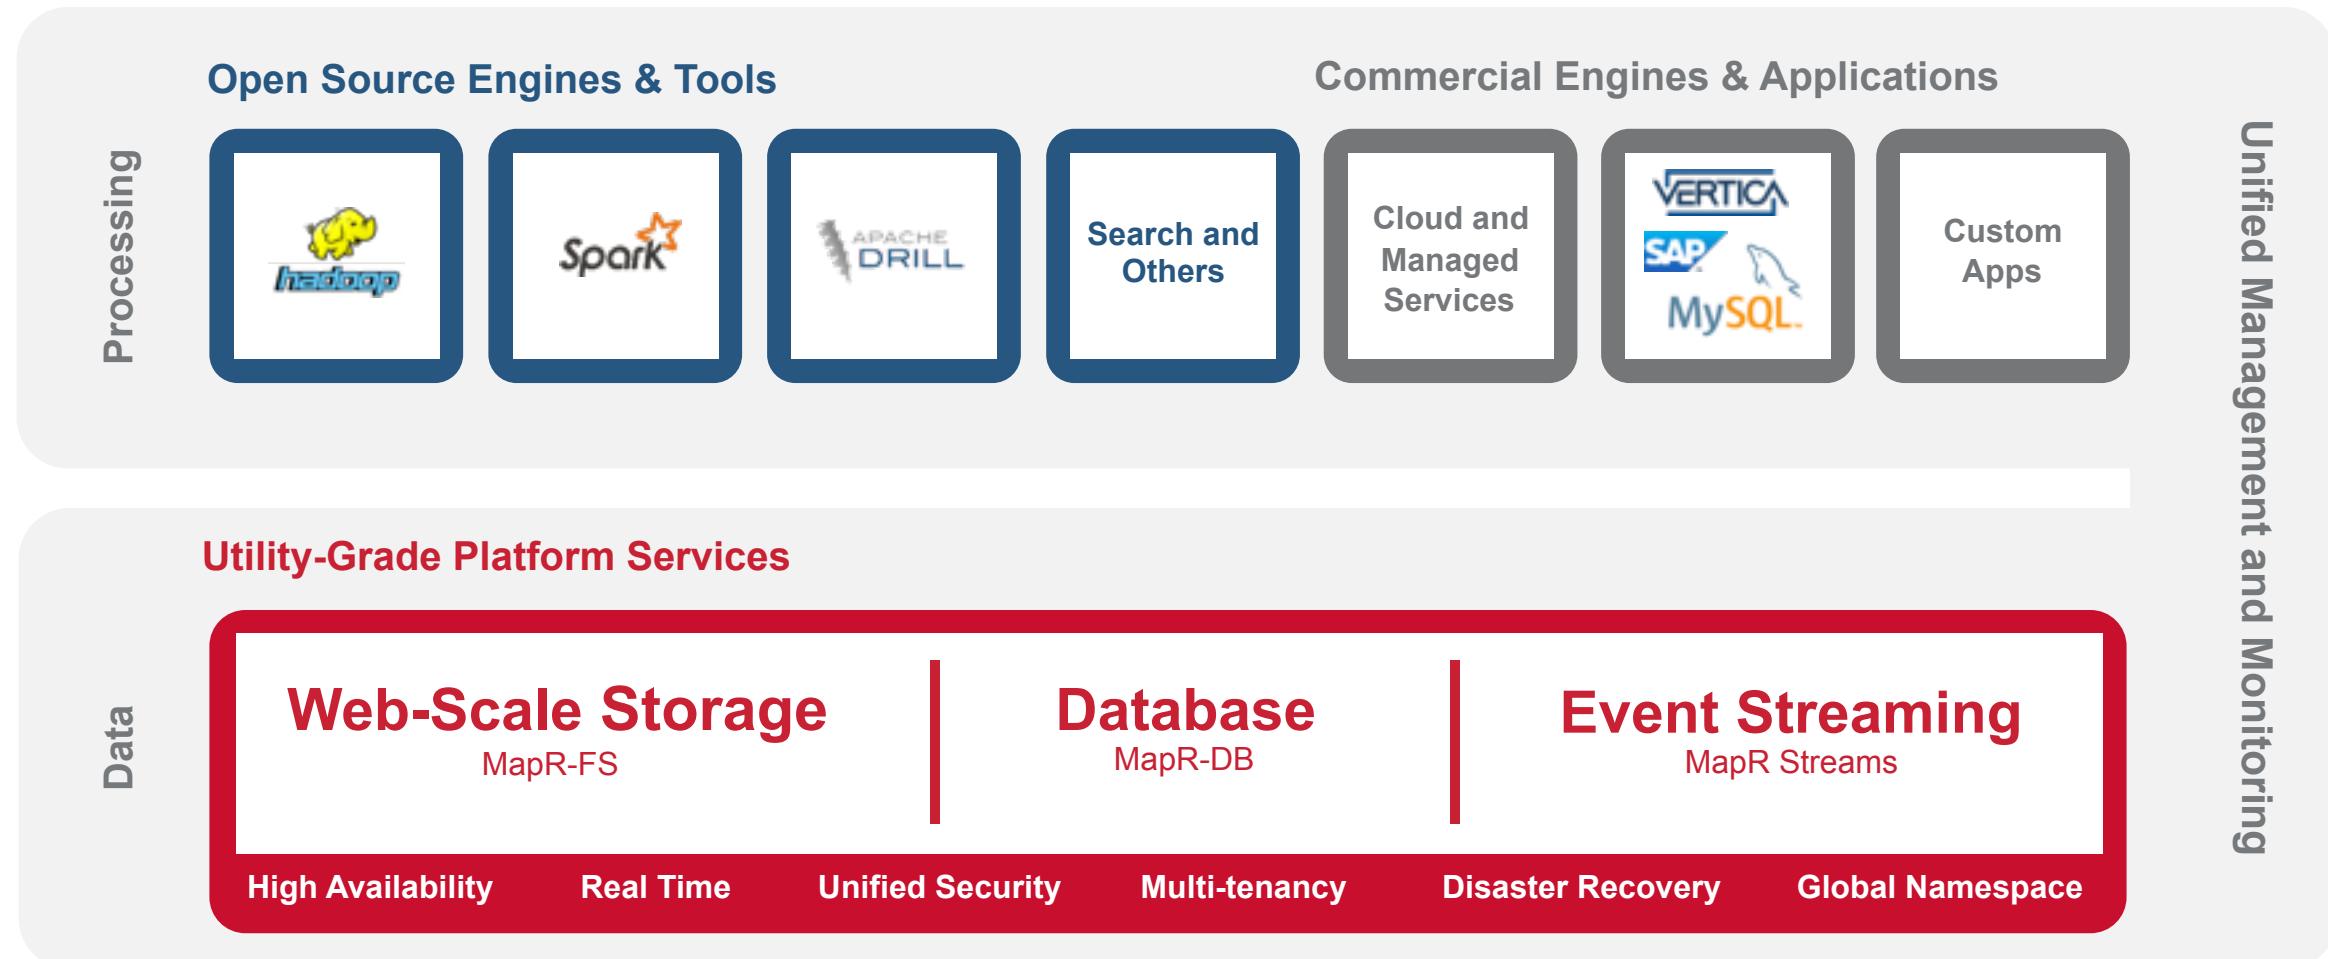


WHAT'S NEXT?

PROBABLY PIZZA-ORIENTED ARCHITECTURE

By @benorama

MapR Converged Data Platform

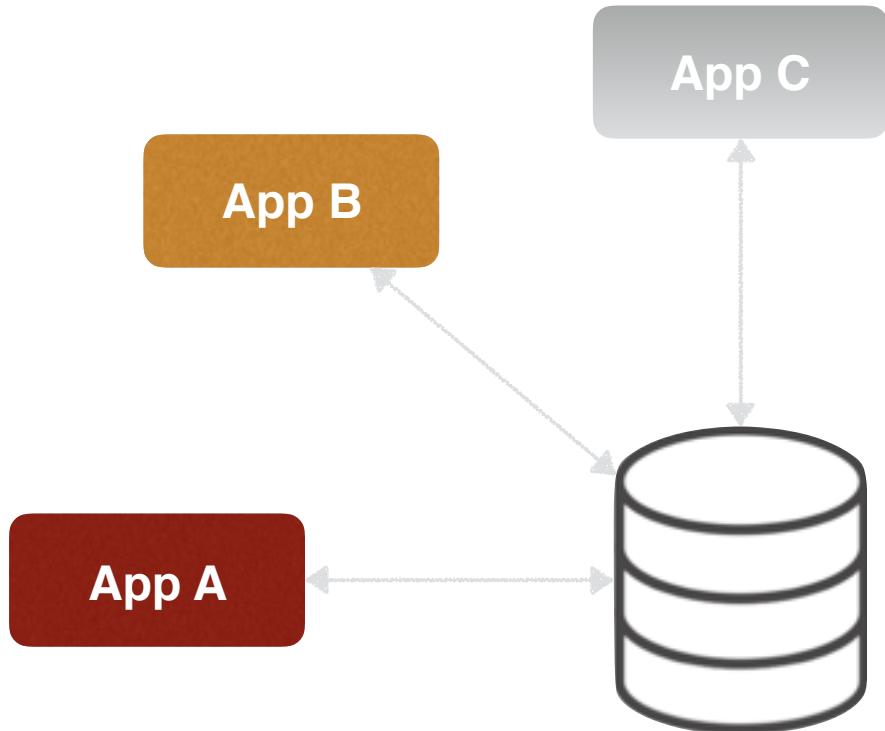


Streaming

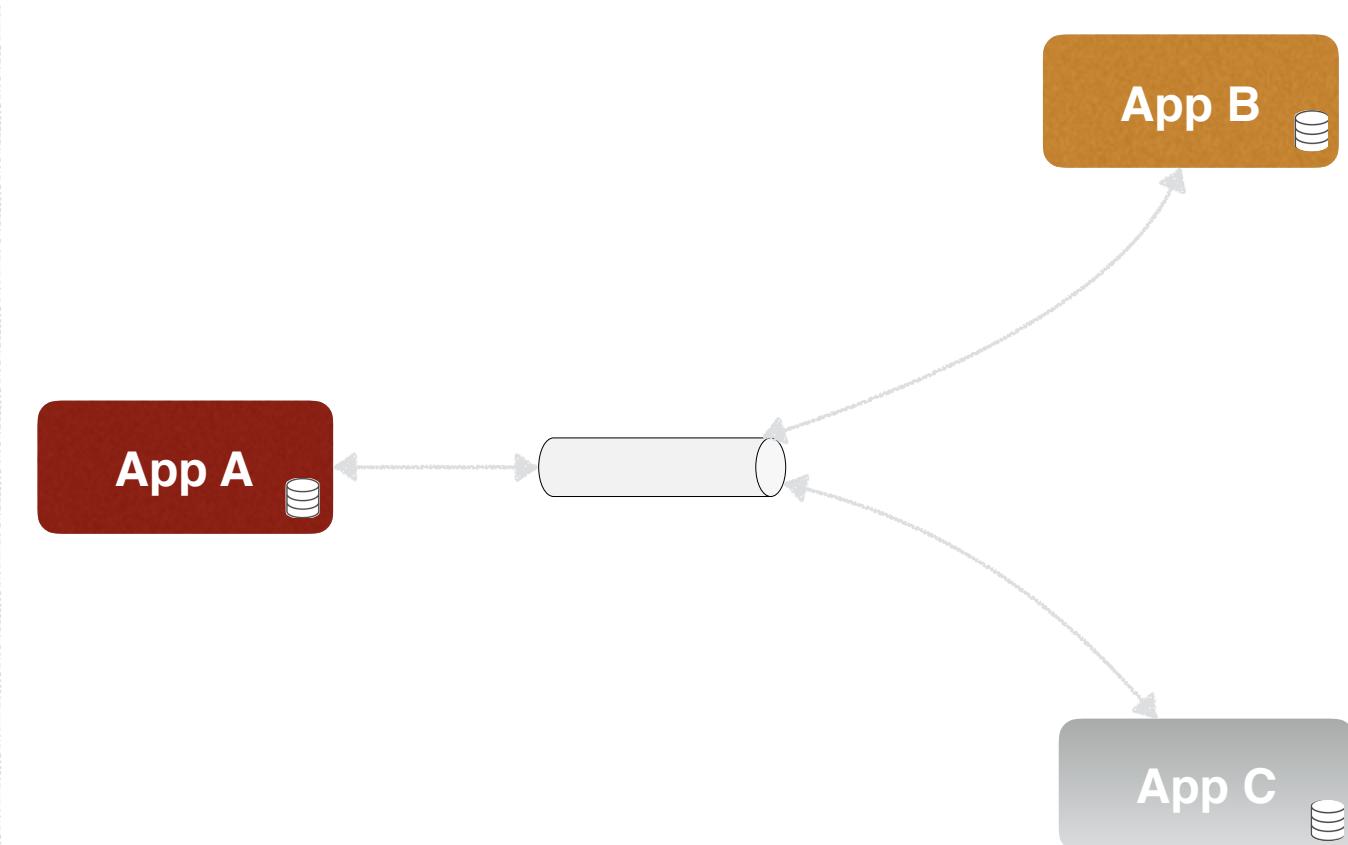
Streaming technology is enabling the obvious:
continuous processing on data
that is **continuously produced**

Hint: you already have streaming data

Decoupling



State managed centralized



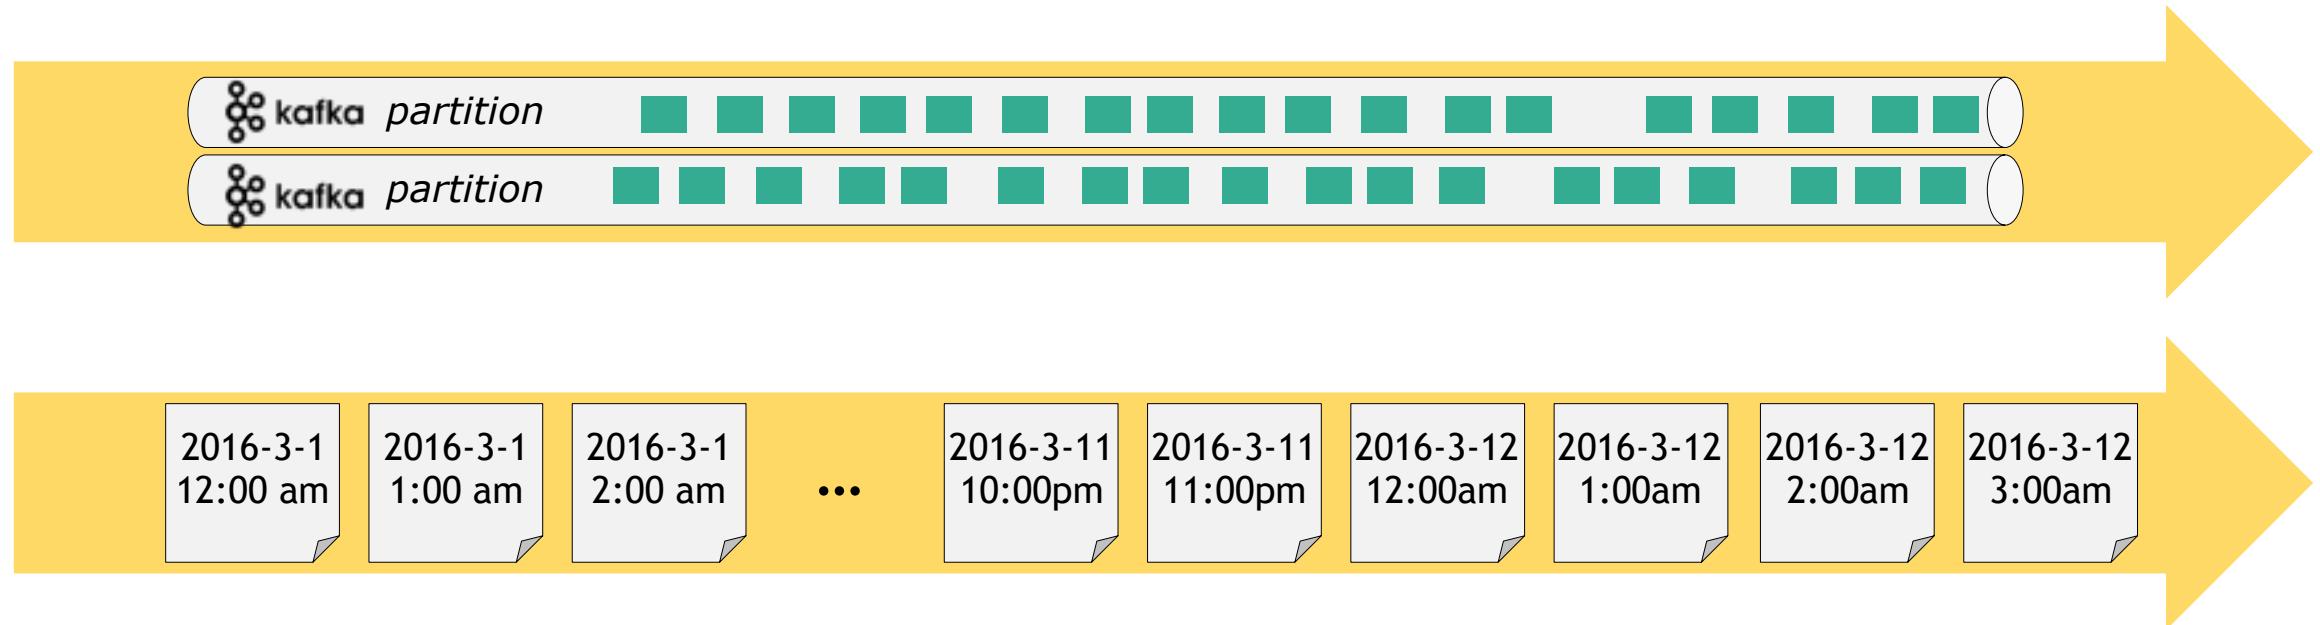
Applications build their own state

Data Pipelines



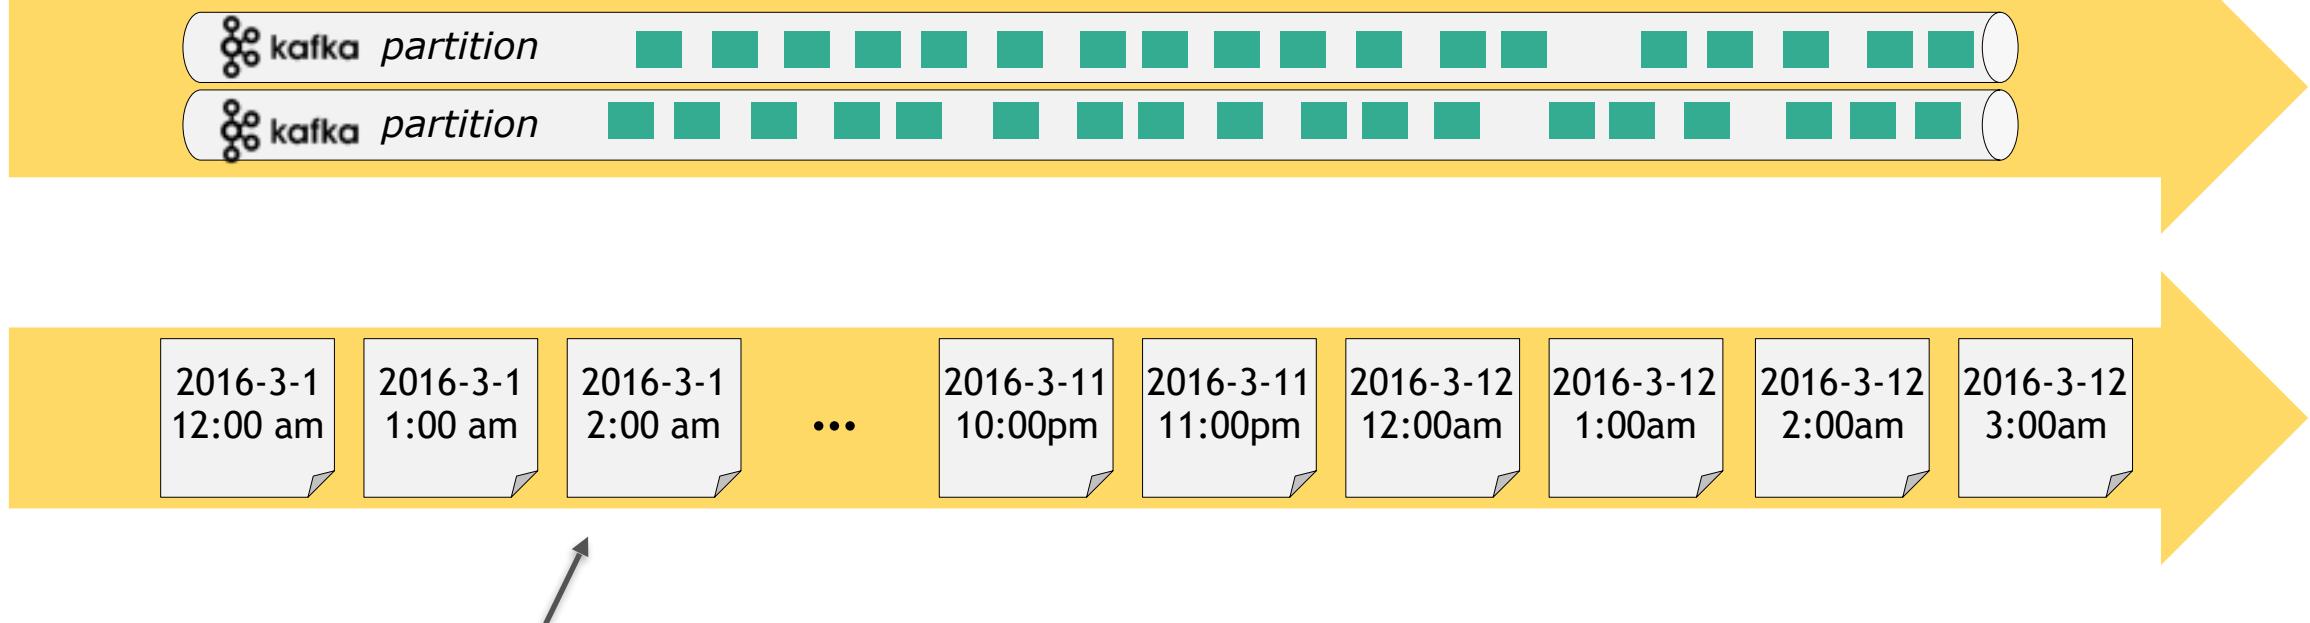
Event Stream

Streaming and Batch



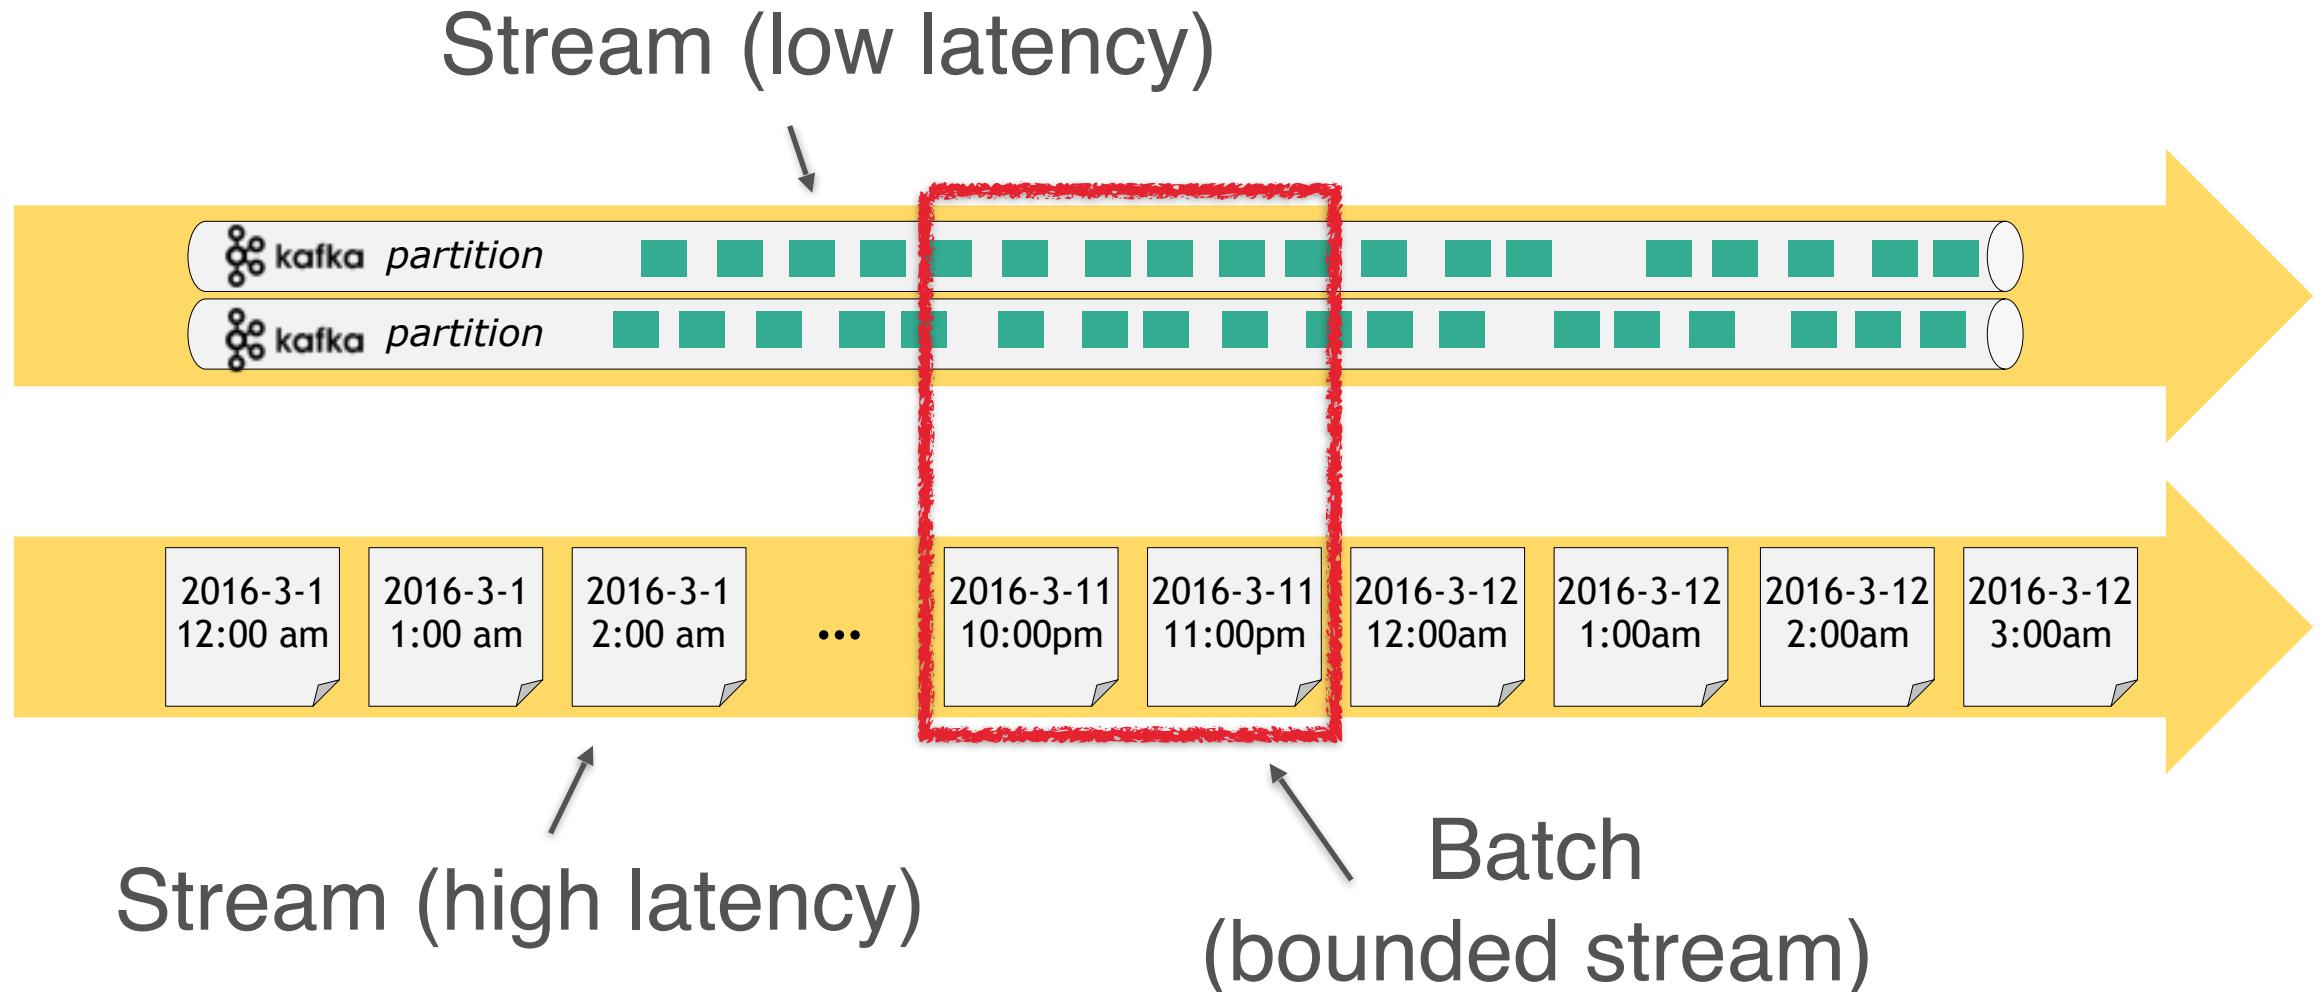
Streaming and Batch

Stream (low latency)



Stream (high latency)

Streaming and Batch



Processing

- Request / Response

Processing

- Request / Response
- Batch

Processing

- Request / Response
- Batch
- Stream Processing

Processing

- Request / Response
- Batch
- Stream Processing
 - Real-time reaction to events
 - Continuous applications
 - Process both real-time and historical data



dataArtisans

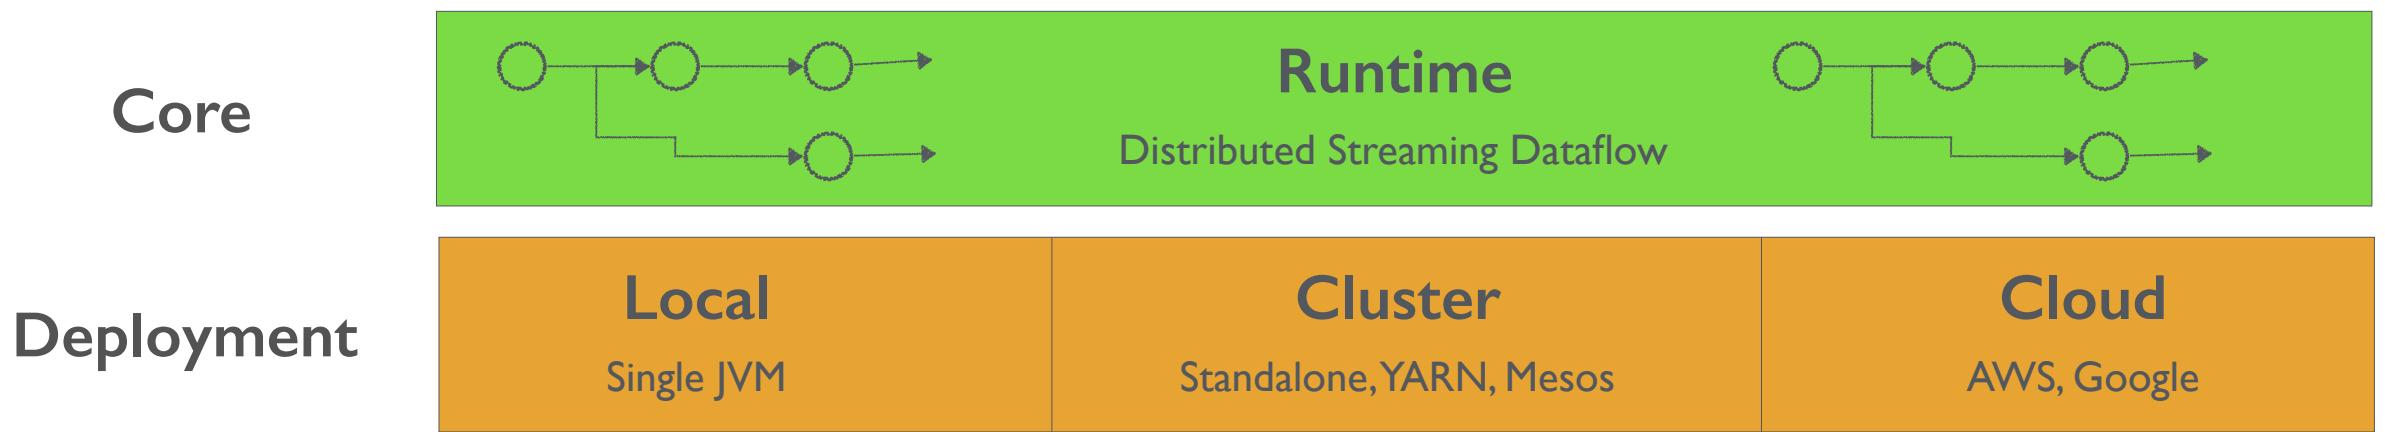
Flink Architecture

Flink Architecture

Deployment

	Local	Cluster	Cloud
	Single JVM	Standalone, YARN, Mesos	AWS, Google

Flink Architecture

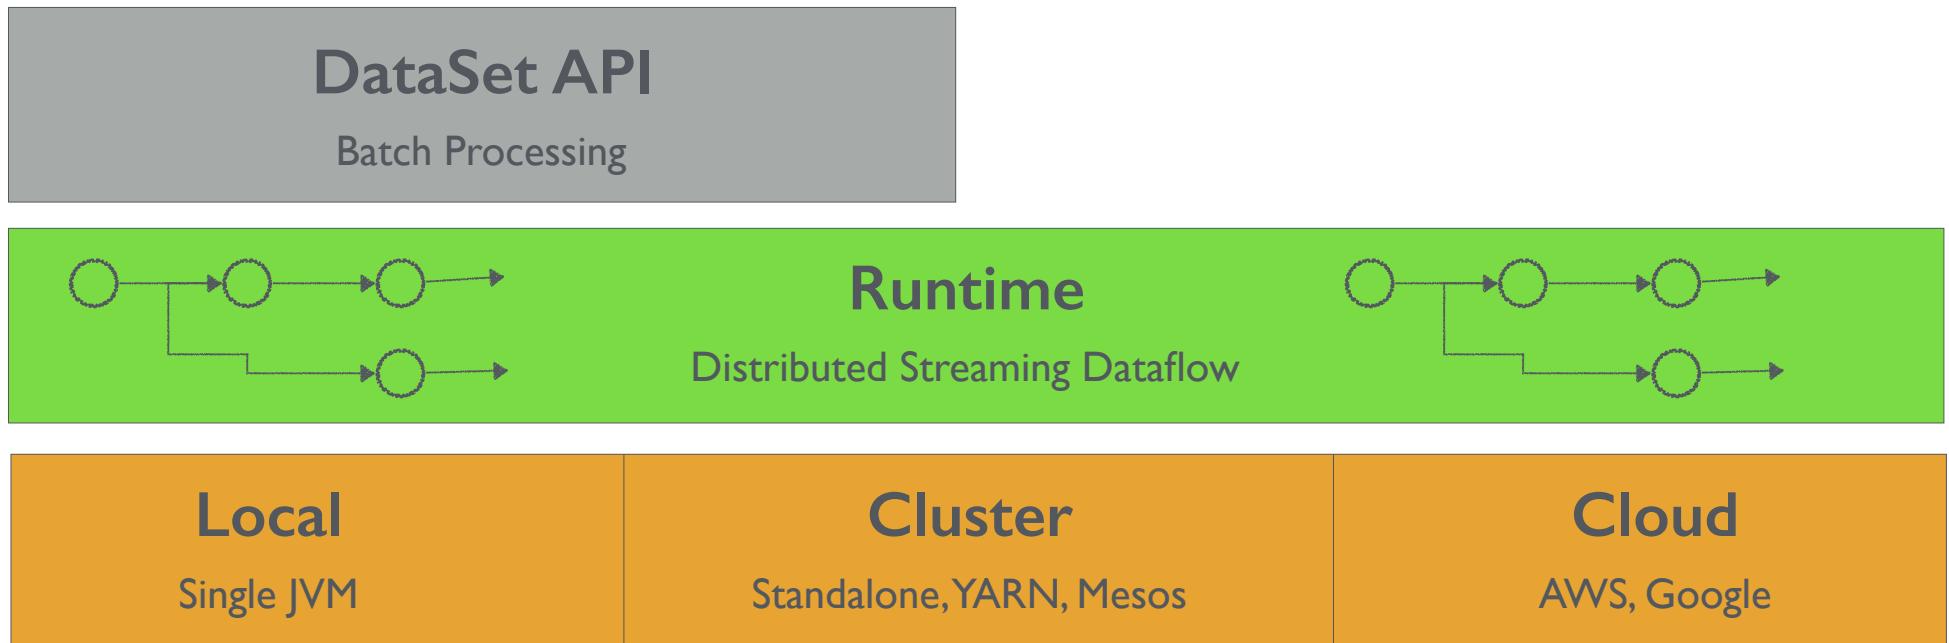


Flink Architecture

API
&
Libraries

Core

Deployment

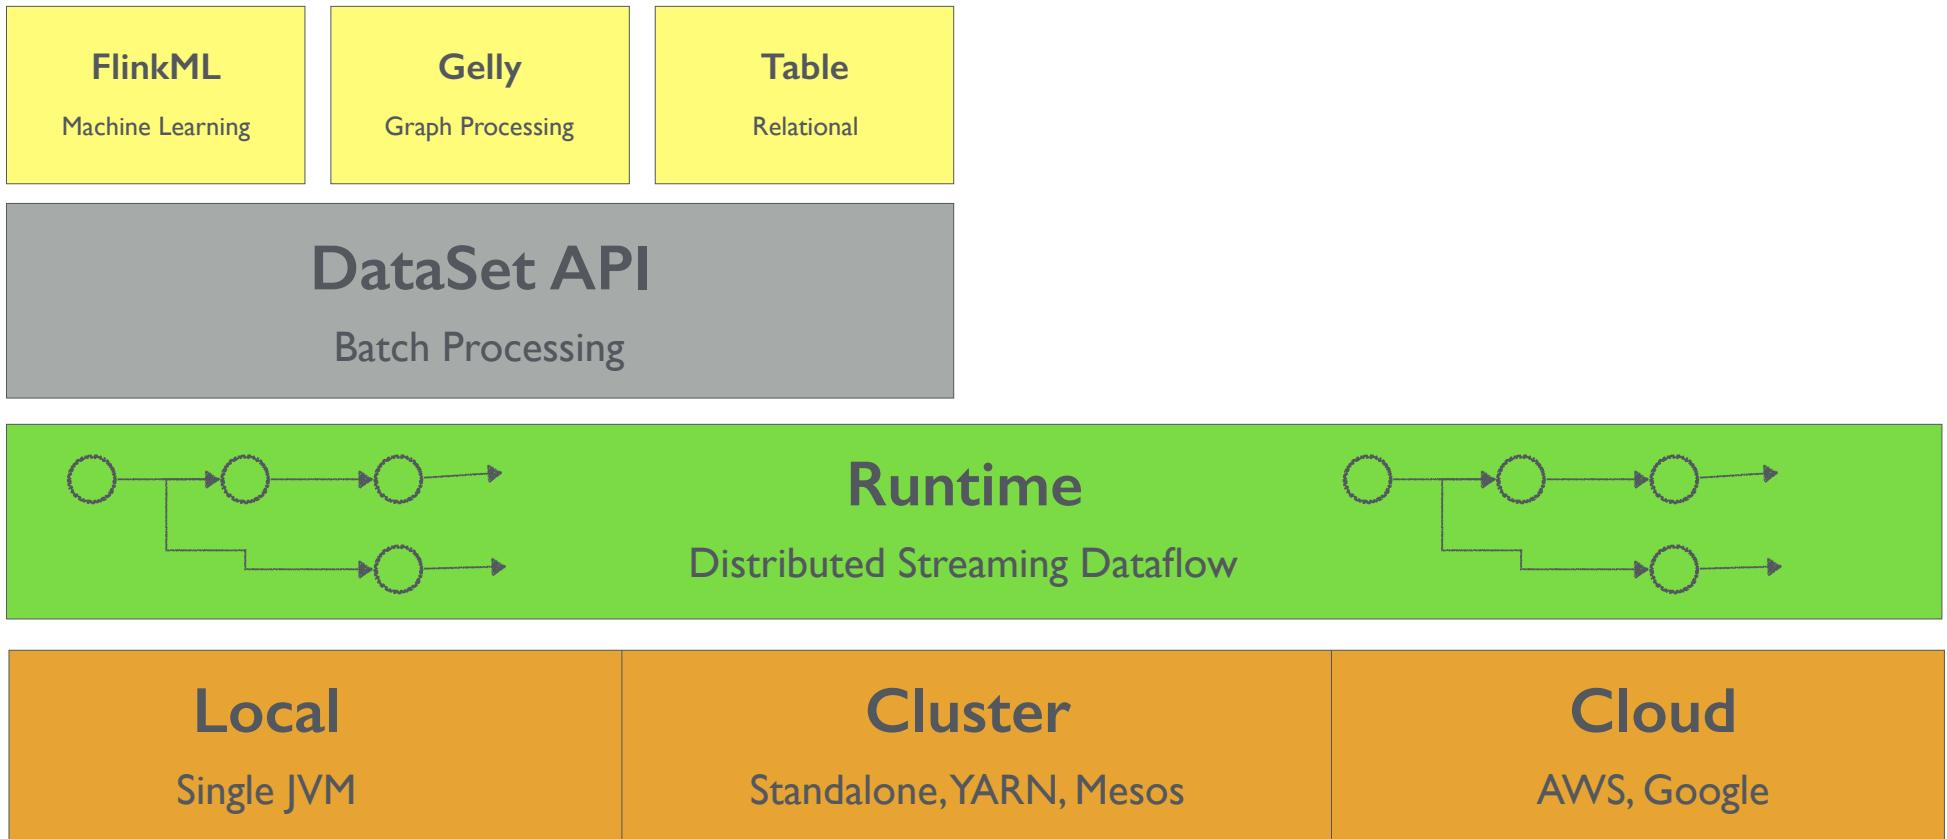


Flink Architecture

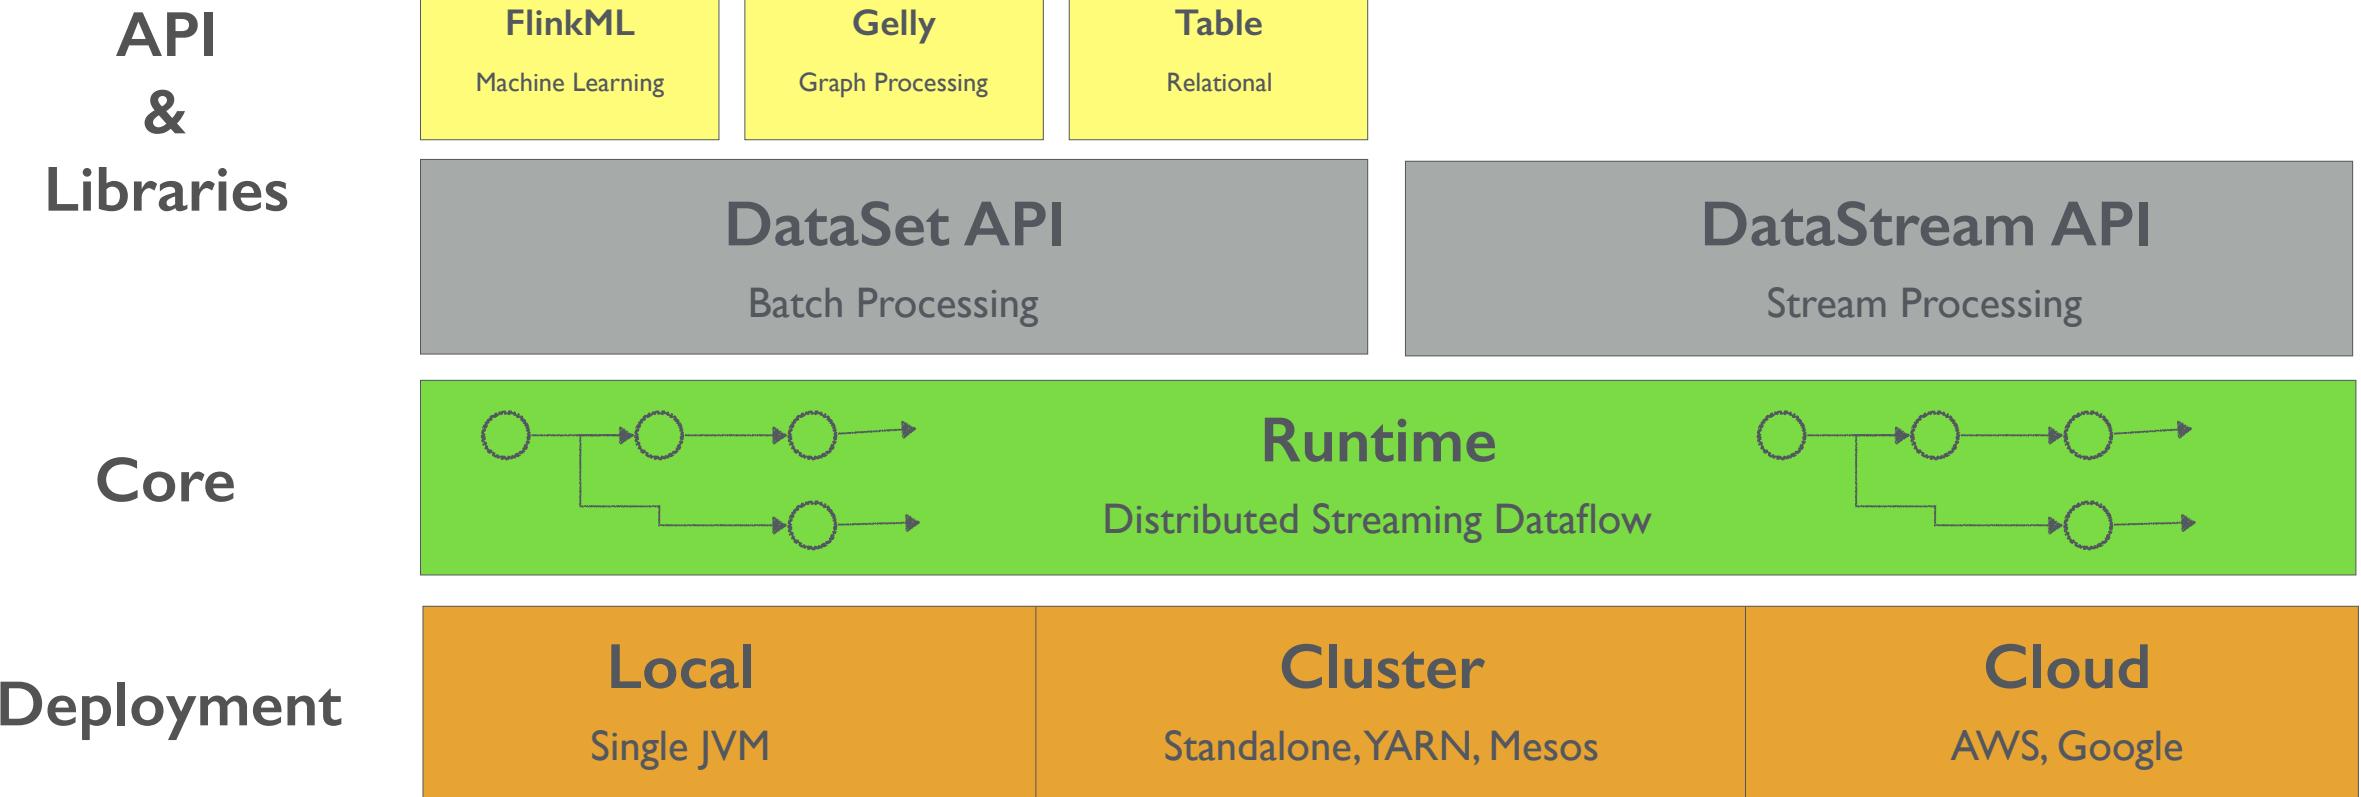
API
&
Libraries

Core

Deployment

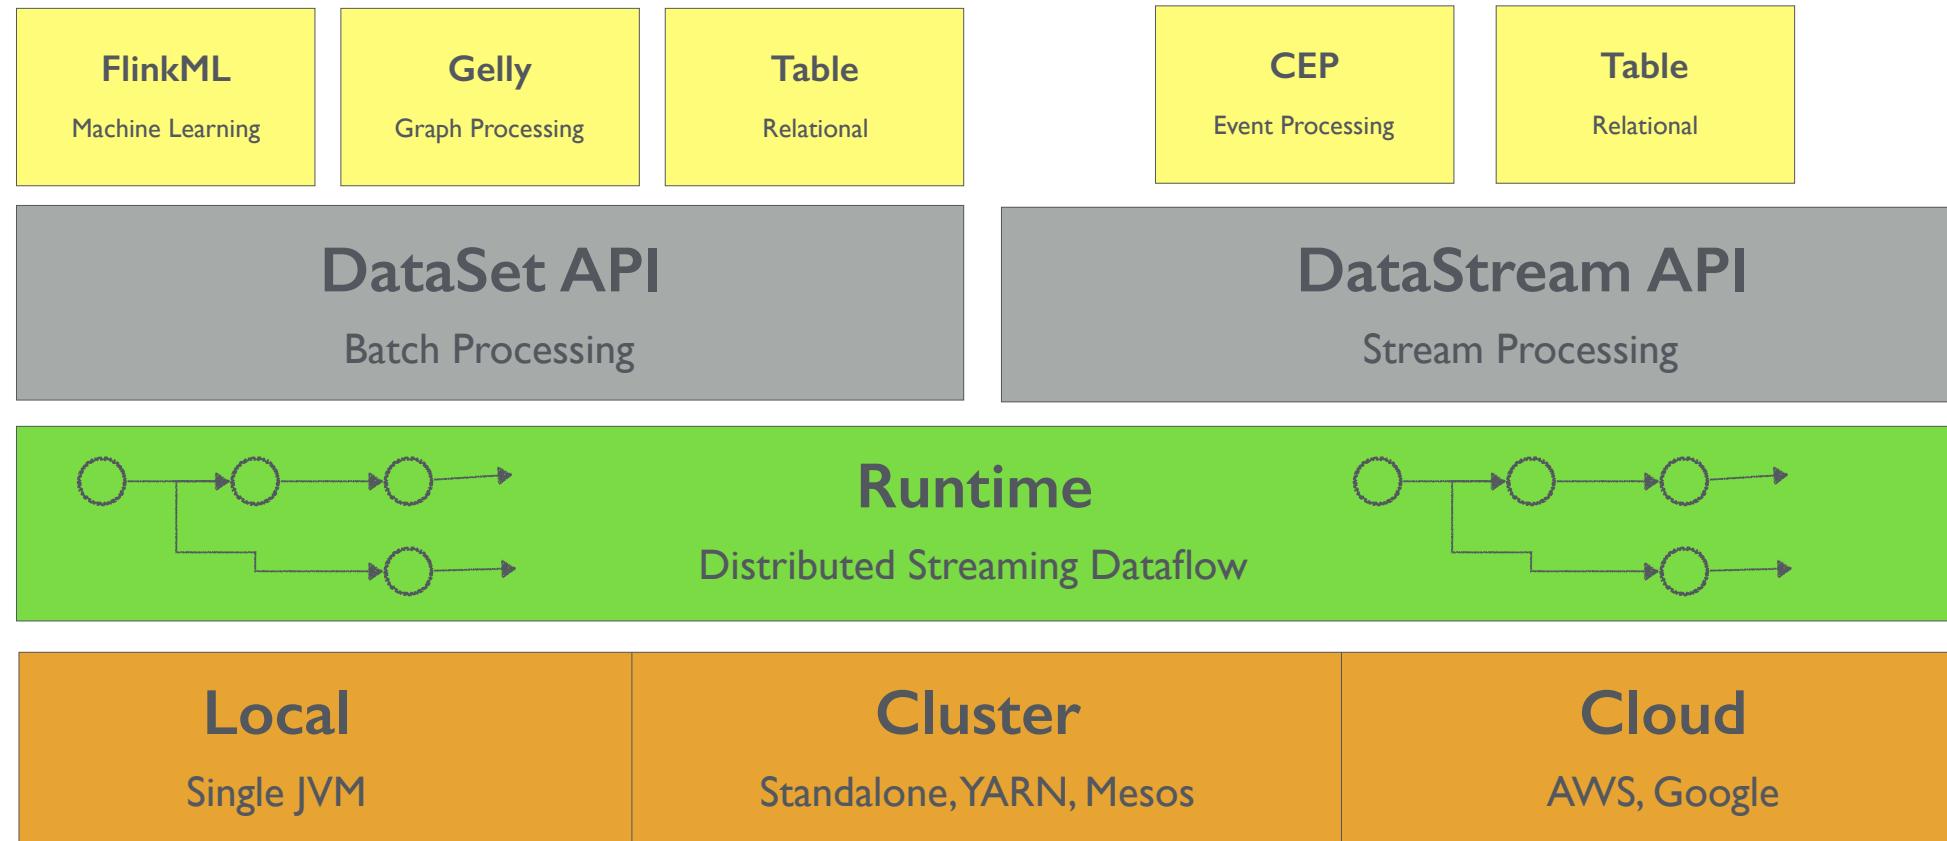


Flink Architecture



Flink Architecture

API
&
Libraries



Demonstration

Flink Basics

Batch & Stream

```
case class Word (word: String, frequency: Int)

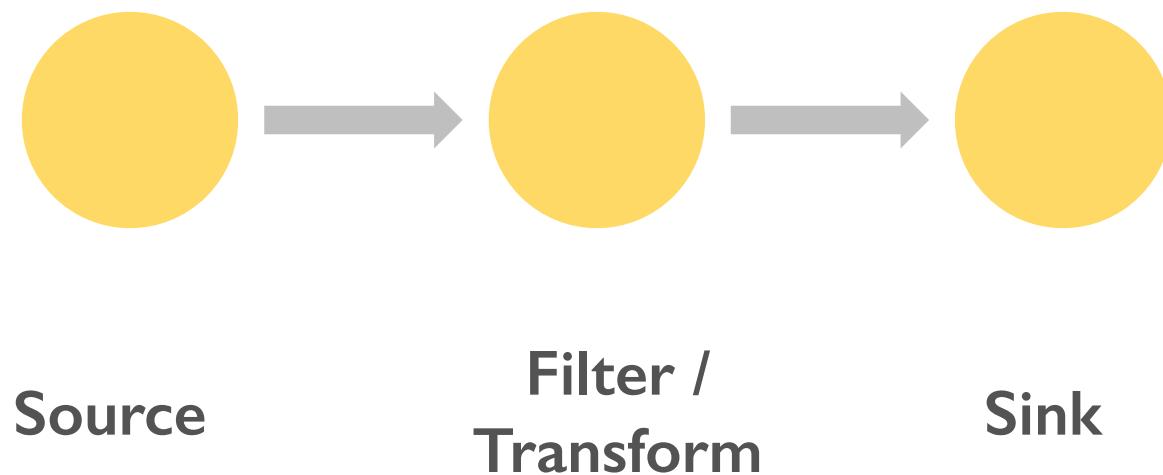
// DataSet API - Batch
val lines: DataSet[String] = env.readTextFile(...)

lines.flatMap {line => line.split(" ")}  
    .map(word => Word(word,1))  
    .groupBy("word").sum("frequency")  
    .print()

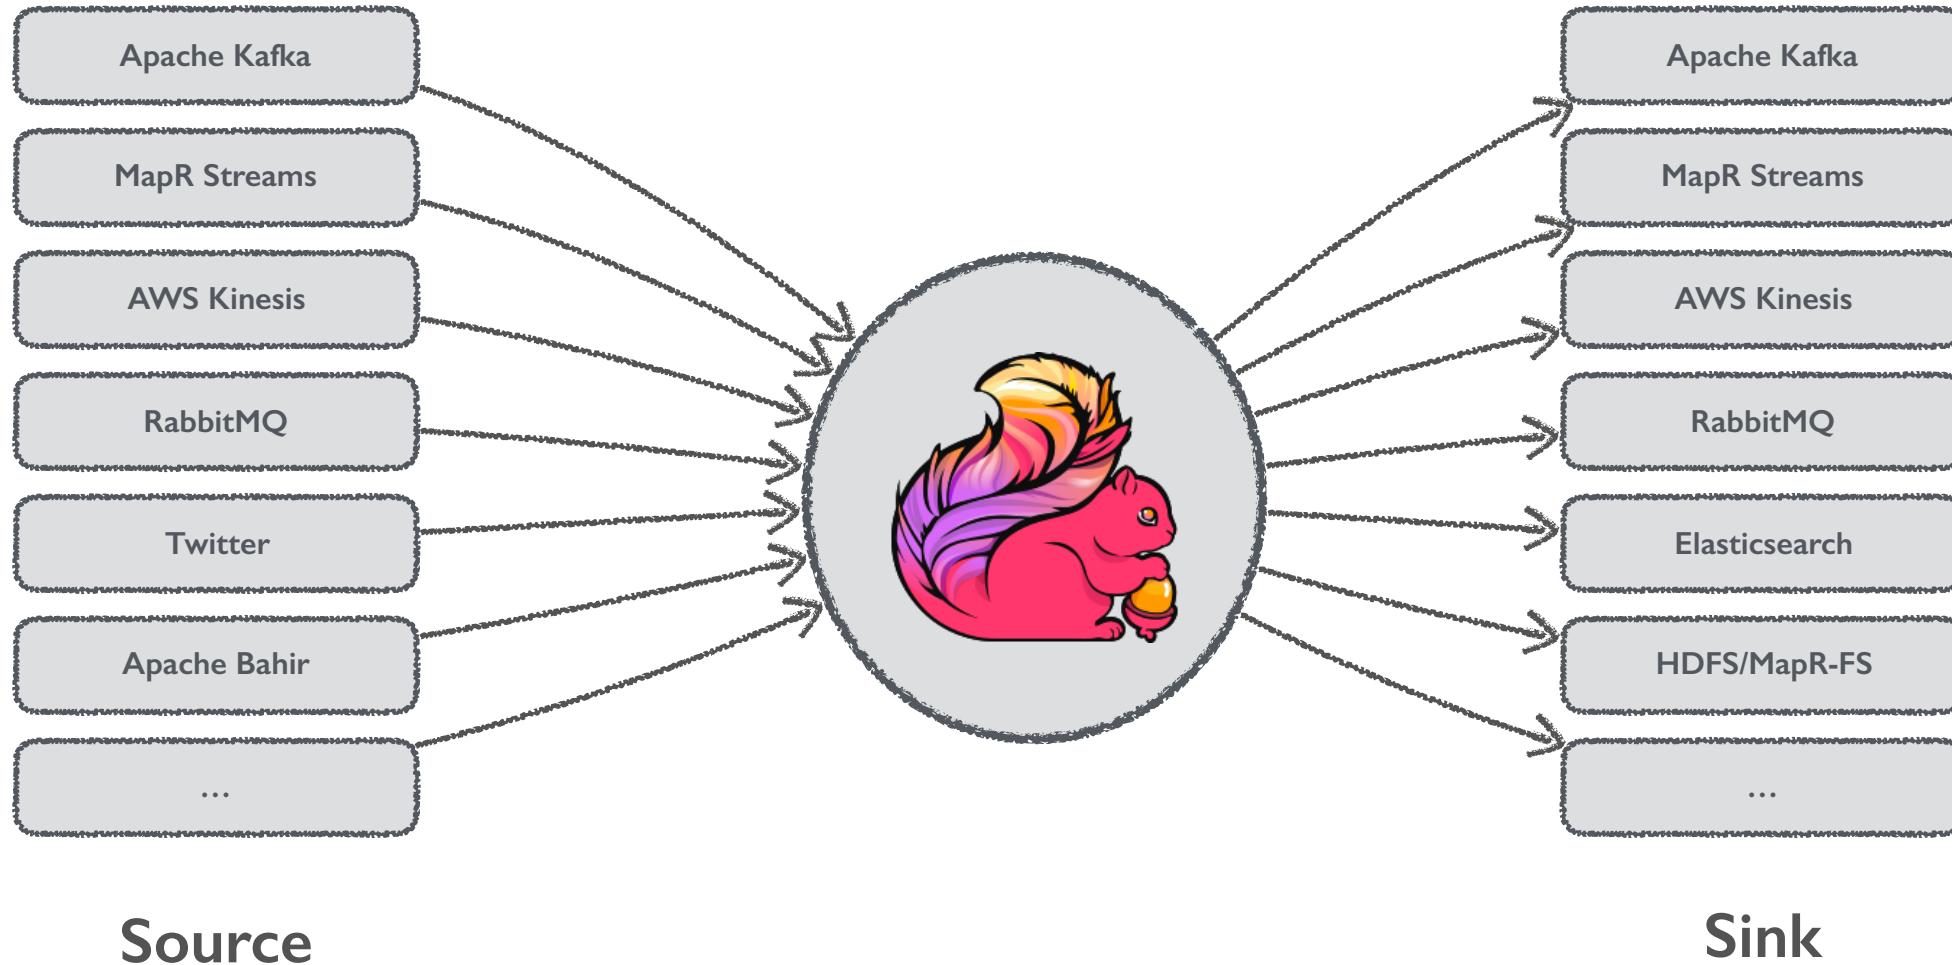
// DataStream API - Streaming
val lines: DataStream[String] = env.fromSocketStream(...)

lines.flatMap {line => line.split(" ")}  
    .map(word => Word(word,1))  
    .keyBy("word").window(Time.of(5,SECONDS))  
    .every(Time.of(1,SECONDS)).sum("frequency")  
    .print()
```

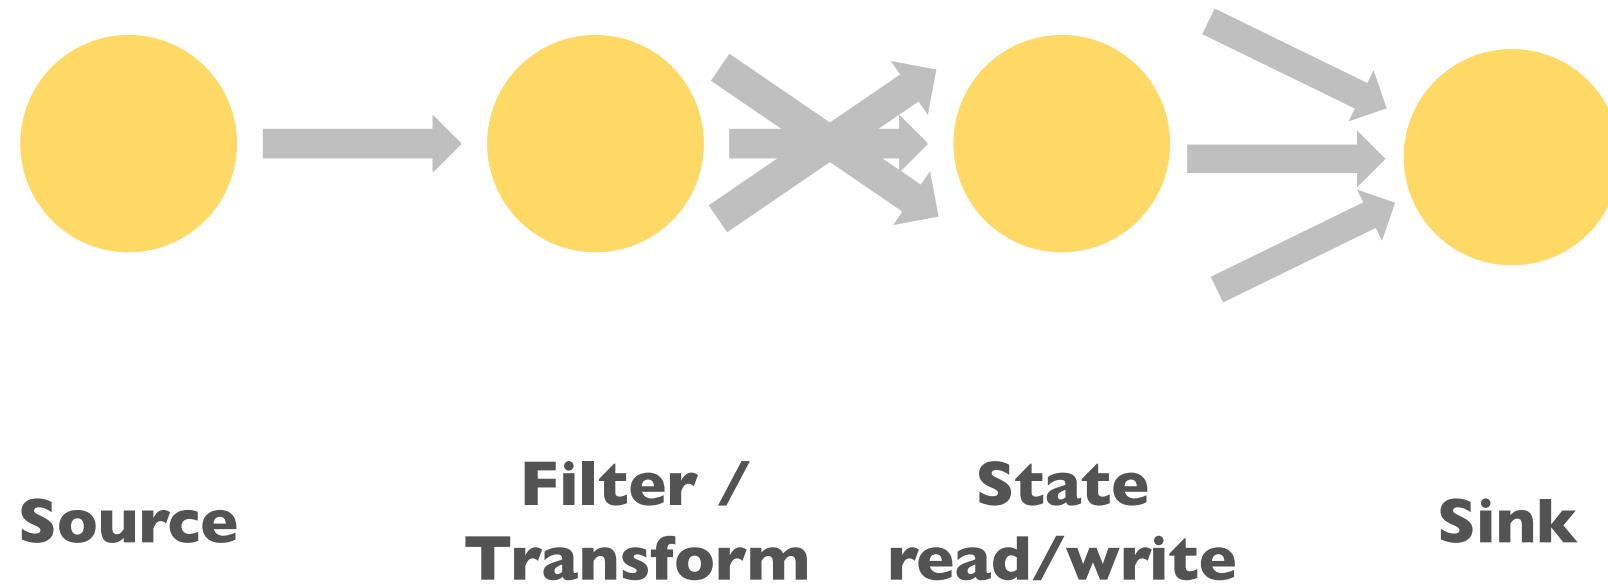
Steam Processing



Flink Ecosystem



Stateful Stream Processing



Is Flink used?

Powered by Flink



otto group

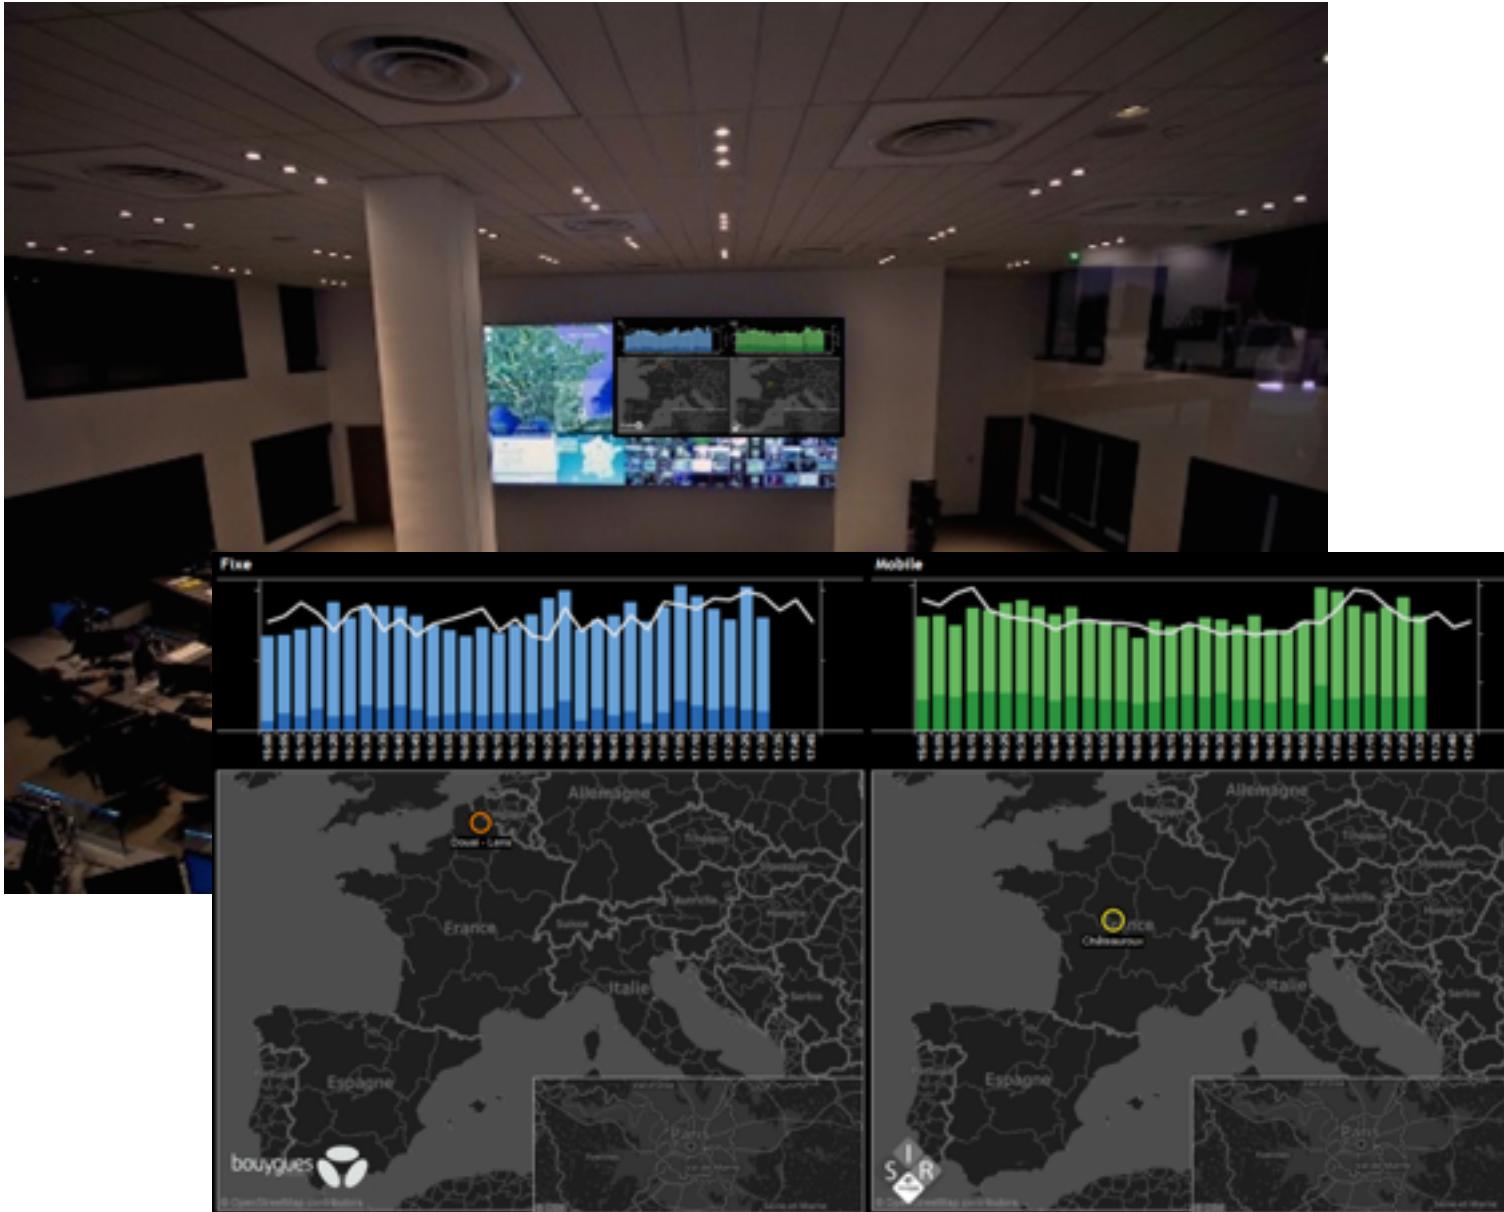
aMADEUS

zalando

Atos

bol.com[€]





10 Billion events/day
2Tb of data/day

30 Applications
2Pb of storage and growing

Source Bouygues Telecom http://berlin.flink-forward.org/wp-content/uploads/2016/07/Thomas-Lamirault_Mohamed-Amine-Abdessemed-A-brief-history-of-time-with-Apache-Flink.pdf

Stream Processing

Windowing

Stream Windows

Sensor → , 9, 6, 8, 4, 7, 3, 8, 4, 2, 1, 3, 2, → out

Stream Windows

Sensor → , 9, 6, 8, 4, 7, 3, 8, 4, 2, 1, 3, 2, →

rolling sum → , 57, 48, 42, 34, 30, 23, 20, 12, 8, 6, 5, 2, → out

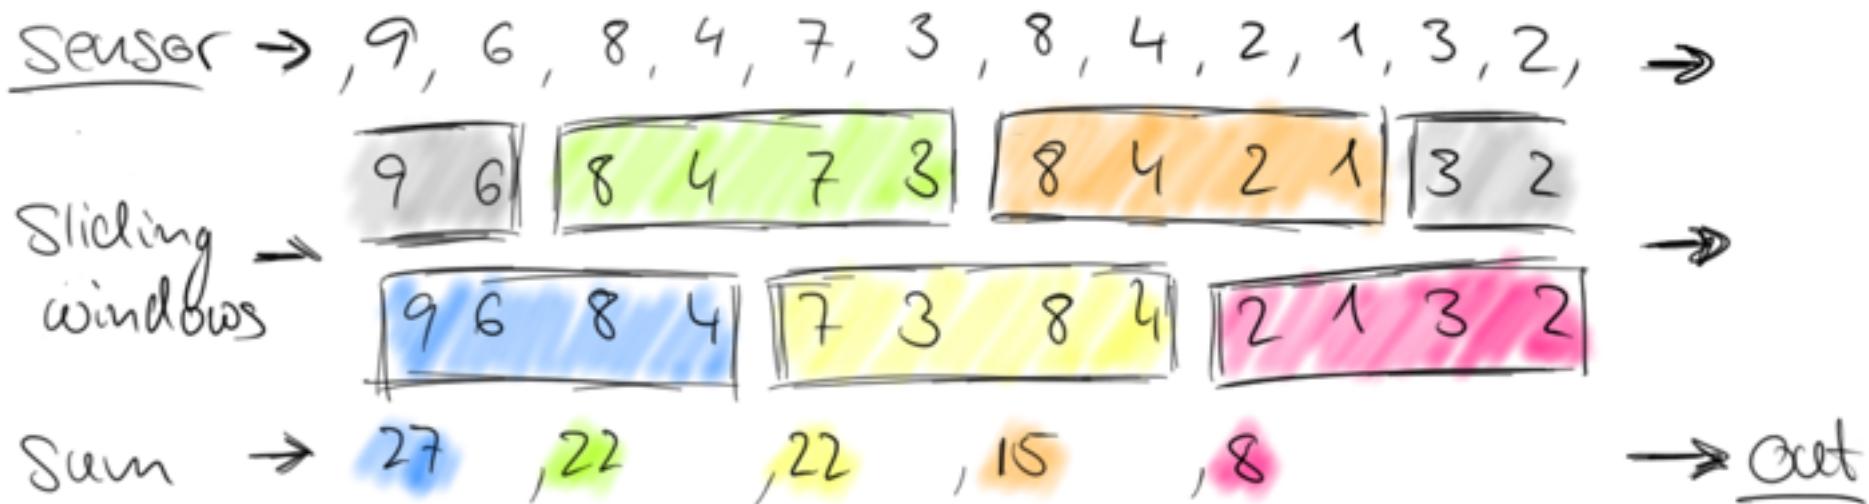
Stream Windows

Sensor $\rightarrow , 9, 6, 8, 4, 7, 3, 8, 4, 2, 1, 3, 2,$ \Rightarrow

tumbling
windows $\rightarrow [9, 6, 8, 4], [7, 3, 8, 4], [2, 1, 3, 2]$ \Rightarrow

sum $\rightarrow 27$, 22 , 8 $\rightarrow \underline{\text{out}}$

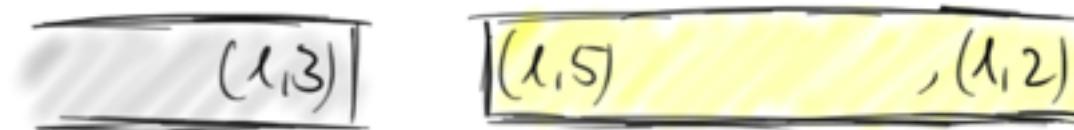
Stream Windows



Stream Windows

Stream \rightarrow $(3,2), (1,3), (2,4), (1,5), (2,1), (3,7), (1,2), (3,4), (3,9), (2,6) \rightarrow$

Tumbling
Count(2) \rightarrow
Windows



Sum \rightarrow $(3,9)$ $(1,7), (2,7)$ $(3,13)$ \rightarrow Out

Demonstration

Flink Windowing

Time

What about it ?

Time in Flink

- Multiple notion of “Time” in Flink
 - Event Time
 - Ingestion Time
 - Processing Time

What Is Event-Time Processing



Event Time

Episode Episode Episode Episode Episode Episode Episode

IV

V

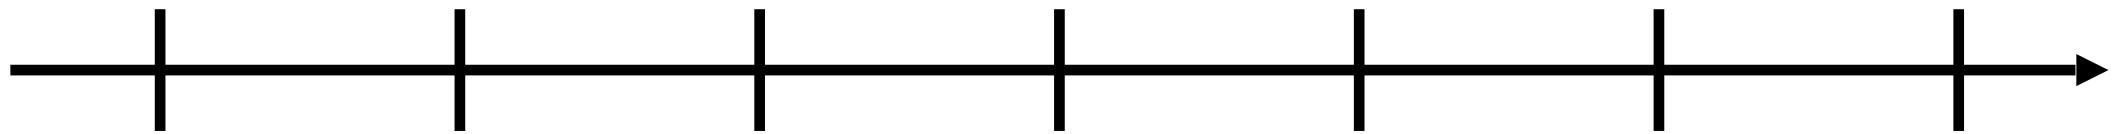
VI

I

II

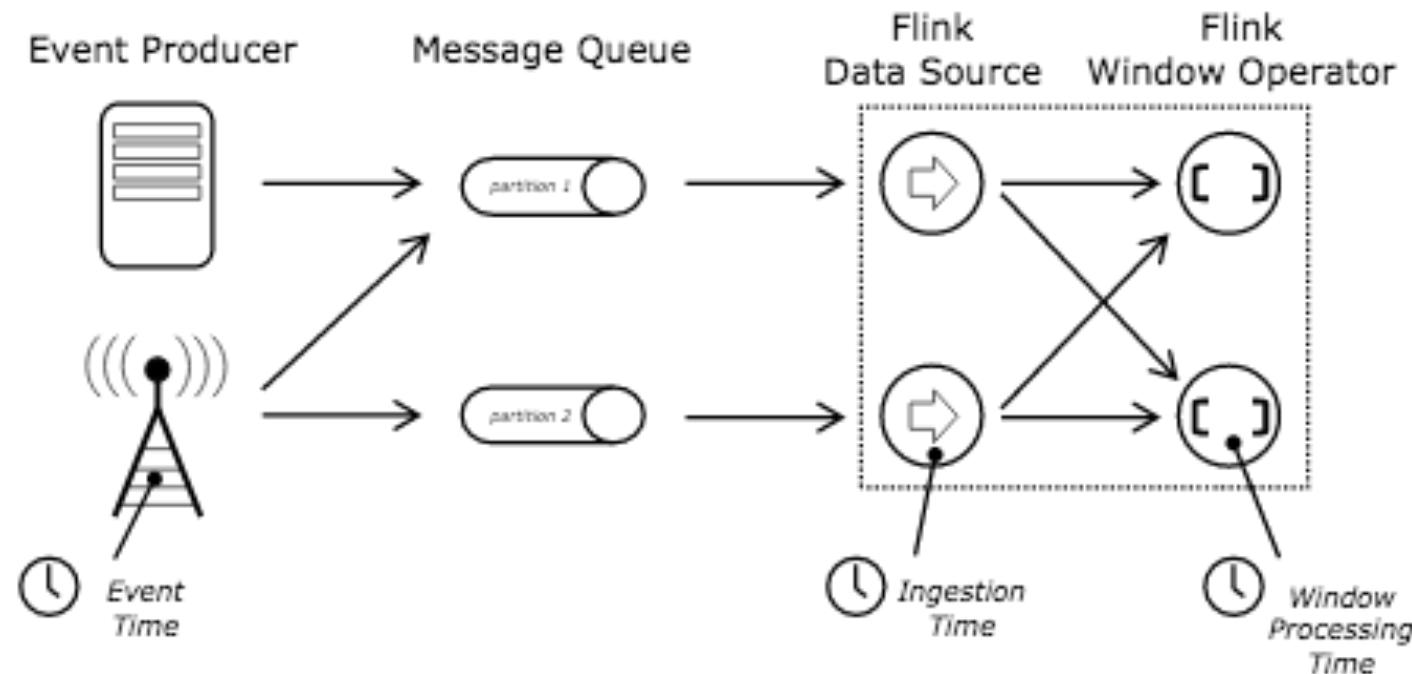
III

VII



Processing Time

Time in Flink

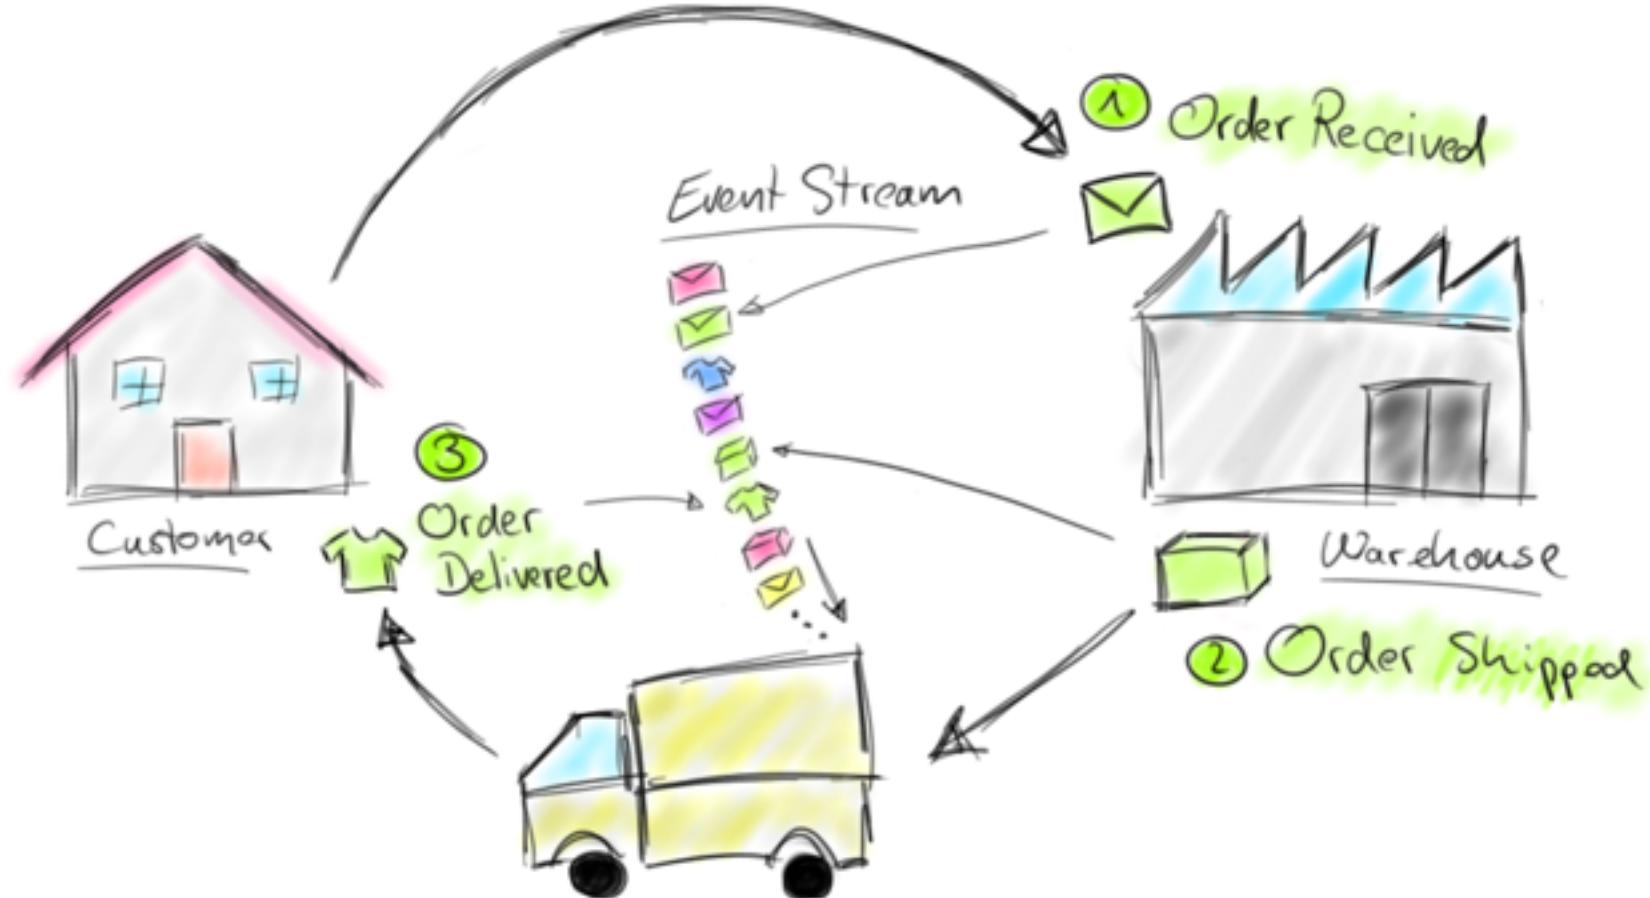


Complex Event Processing

Complex Event Processing

- Analyzing a stream of events and drawing conclusions
 - “if A and then B → infer event C”
- Demanding requirements on stream processor
 - Low latency!
 - Exactly-once semantics & event-time support

Use Case



Order Events

Process is reflected in a stream of order events

Order(orderId, tStamp, "received")

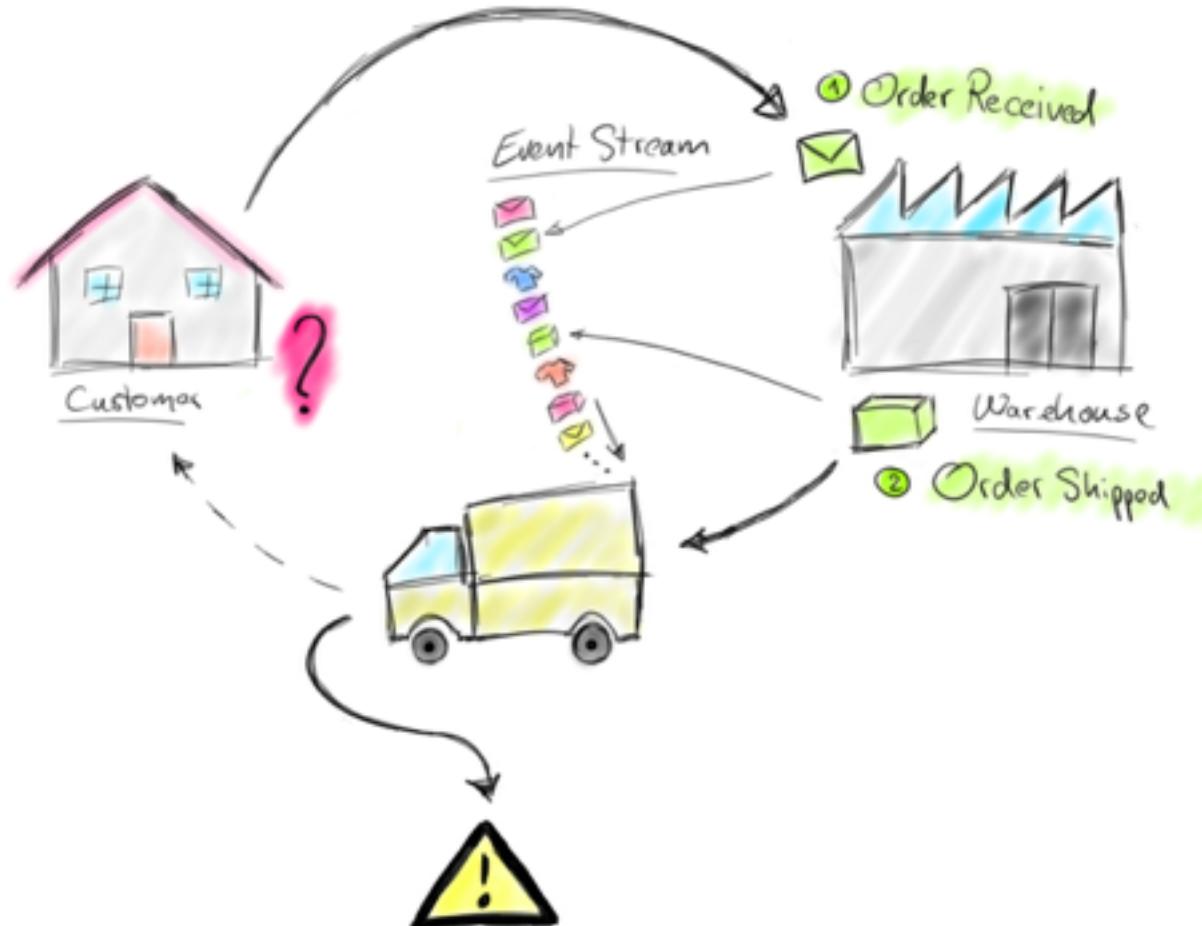
Shipment(orderId, tStamp, "shipped")

Delivery(orderId, tStamp, "delivered")

orderId: Identifies the order

tStamp: Time at which the event happened

Real-time Warnings



CEP to the Rescue

Define processing and delivery intervals (SLAs)

`ProcessSucc(orderId, tStamp, duration)`

`ProcessWarn(orderId, tStamp)`

`DeliverySucc(orderId, tStamp, duration)`

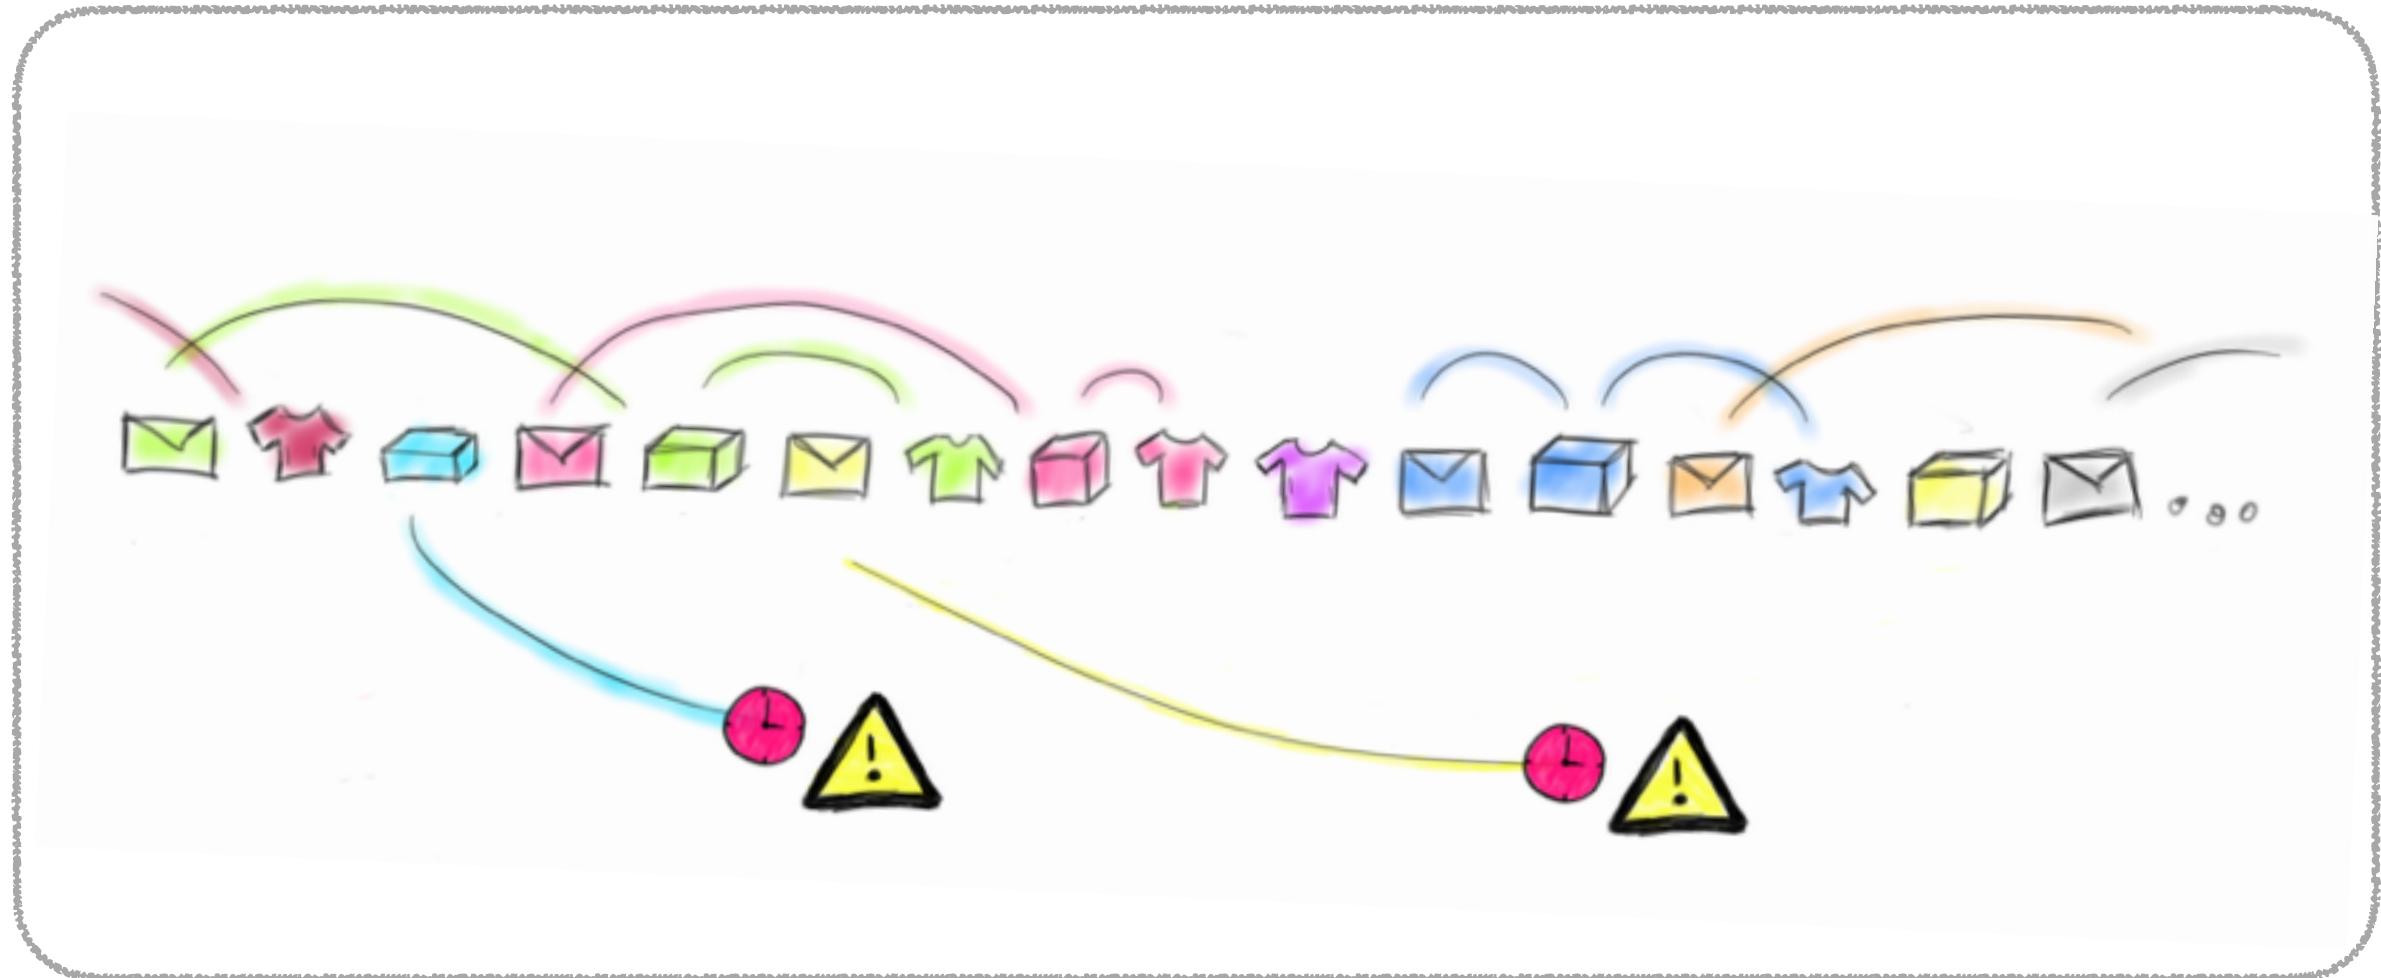
`DeliveryWarn(orderId, tStamp)`

`orderId`: Identifies the order

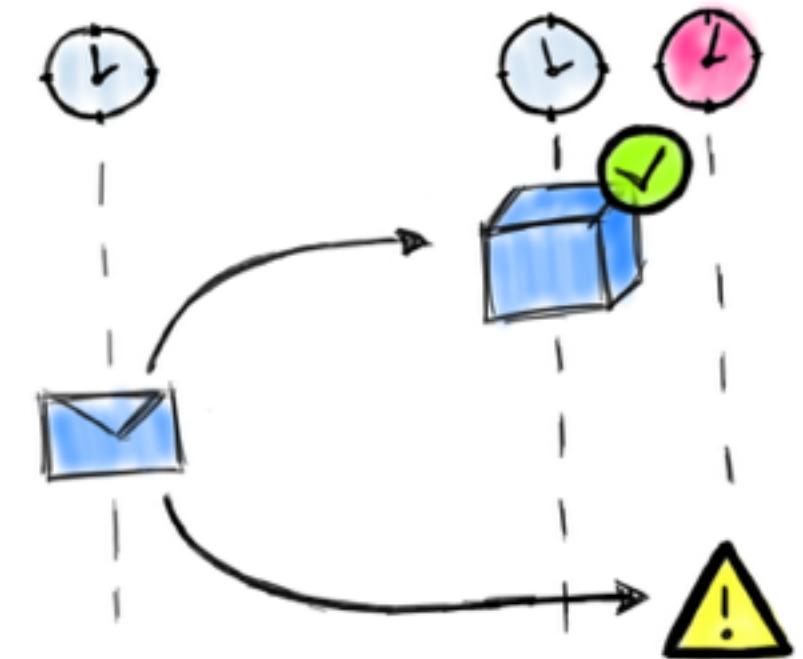
`tStamp`: Time when the event happened

`duration`: Duration of the processing/delivery

CEP Example

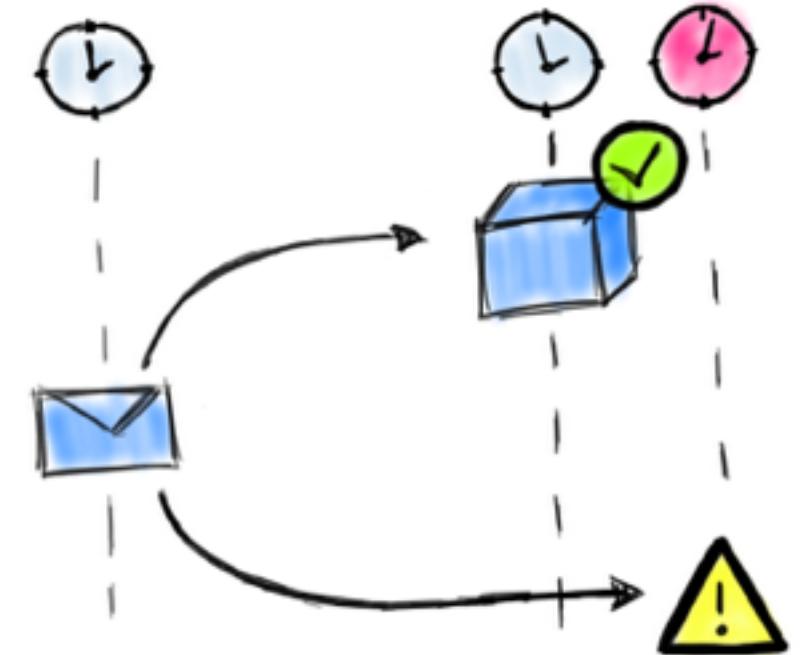


Processing: Order → Shipment



Processing: Order → Shipment

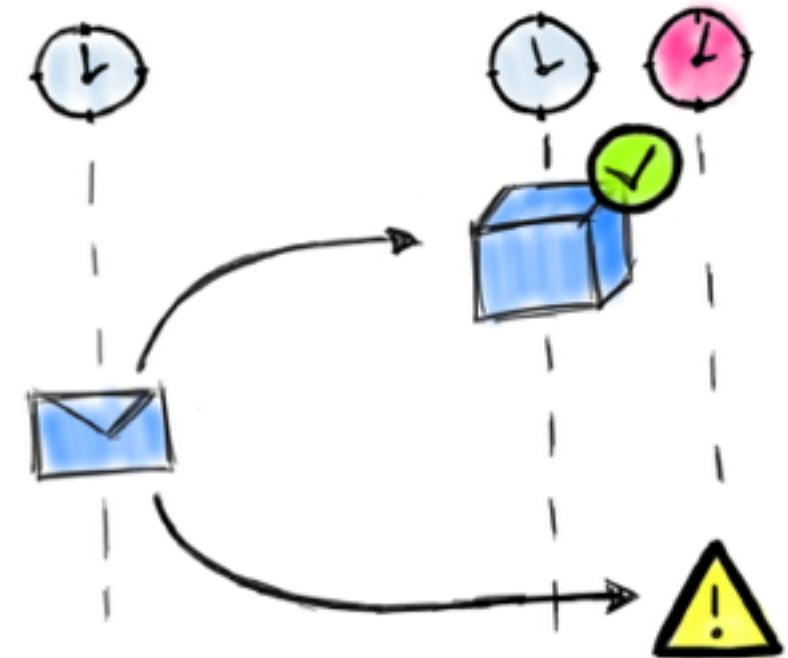
```
val processingPattern = Pattern
  .begin[Event]("received").subtype(classOf[Order])
  .followedBy("shipped").where(_.status == "shipped")
  .within(Time.hours(1))
```



Processing: Order → Shipment

```
val processingPattern = Pattern
    .begin[Event]("received").subtype(classOf[Order])
    .followedBy("shipped").where(_.status == "shipped")
    .within(Time.hours(1))

val processingPatternStream = CEP.pattern(
    input.keyBy("orderId"),
    processingPattern)
```

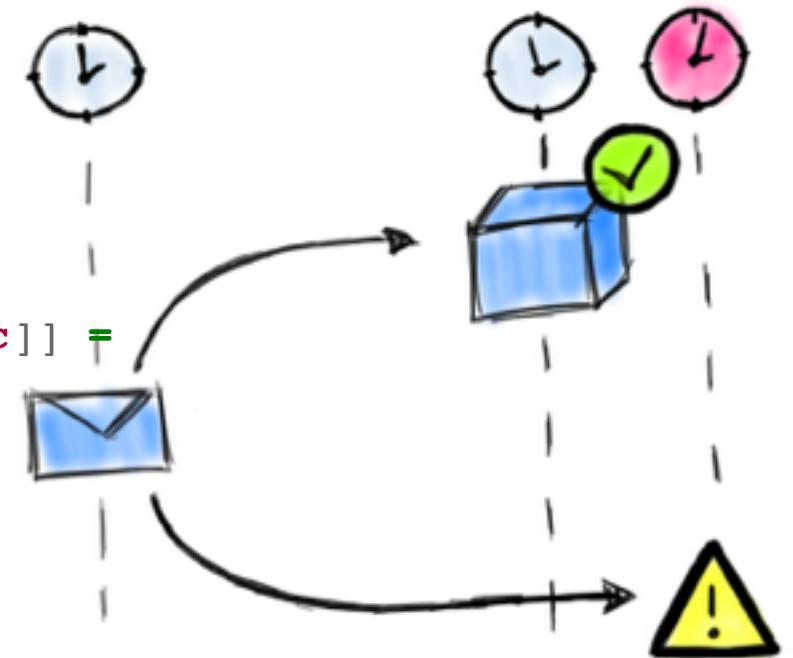


Processing: Order → Shipment

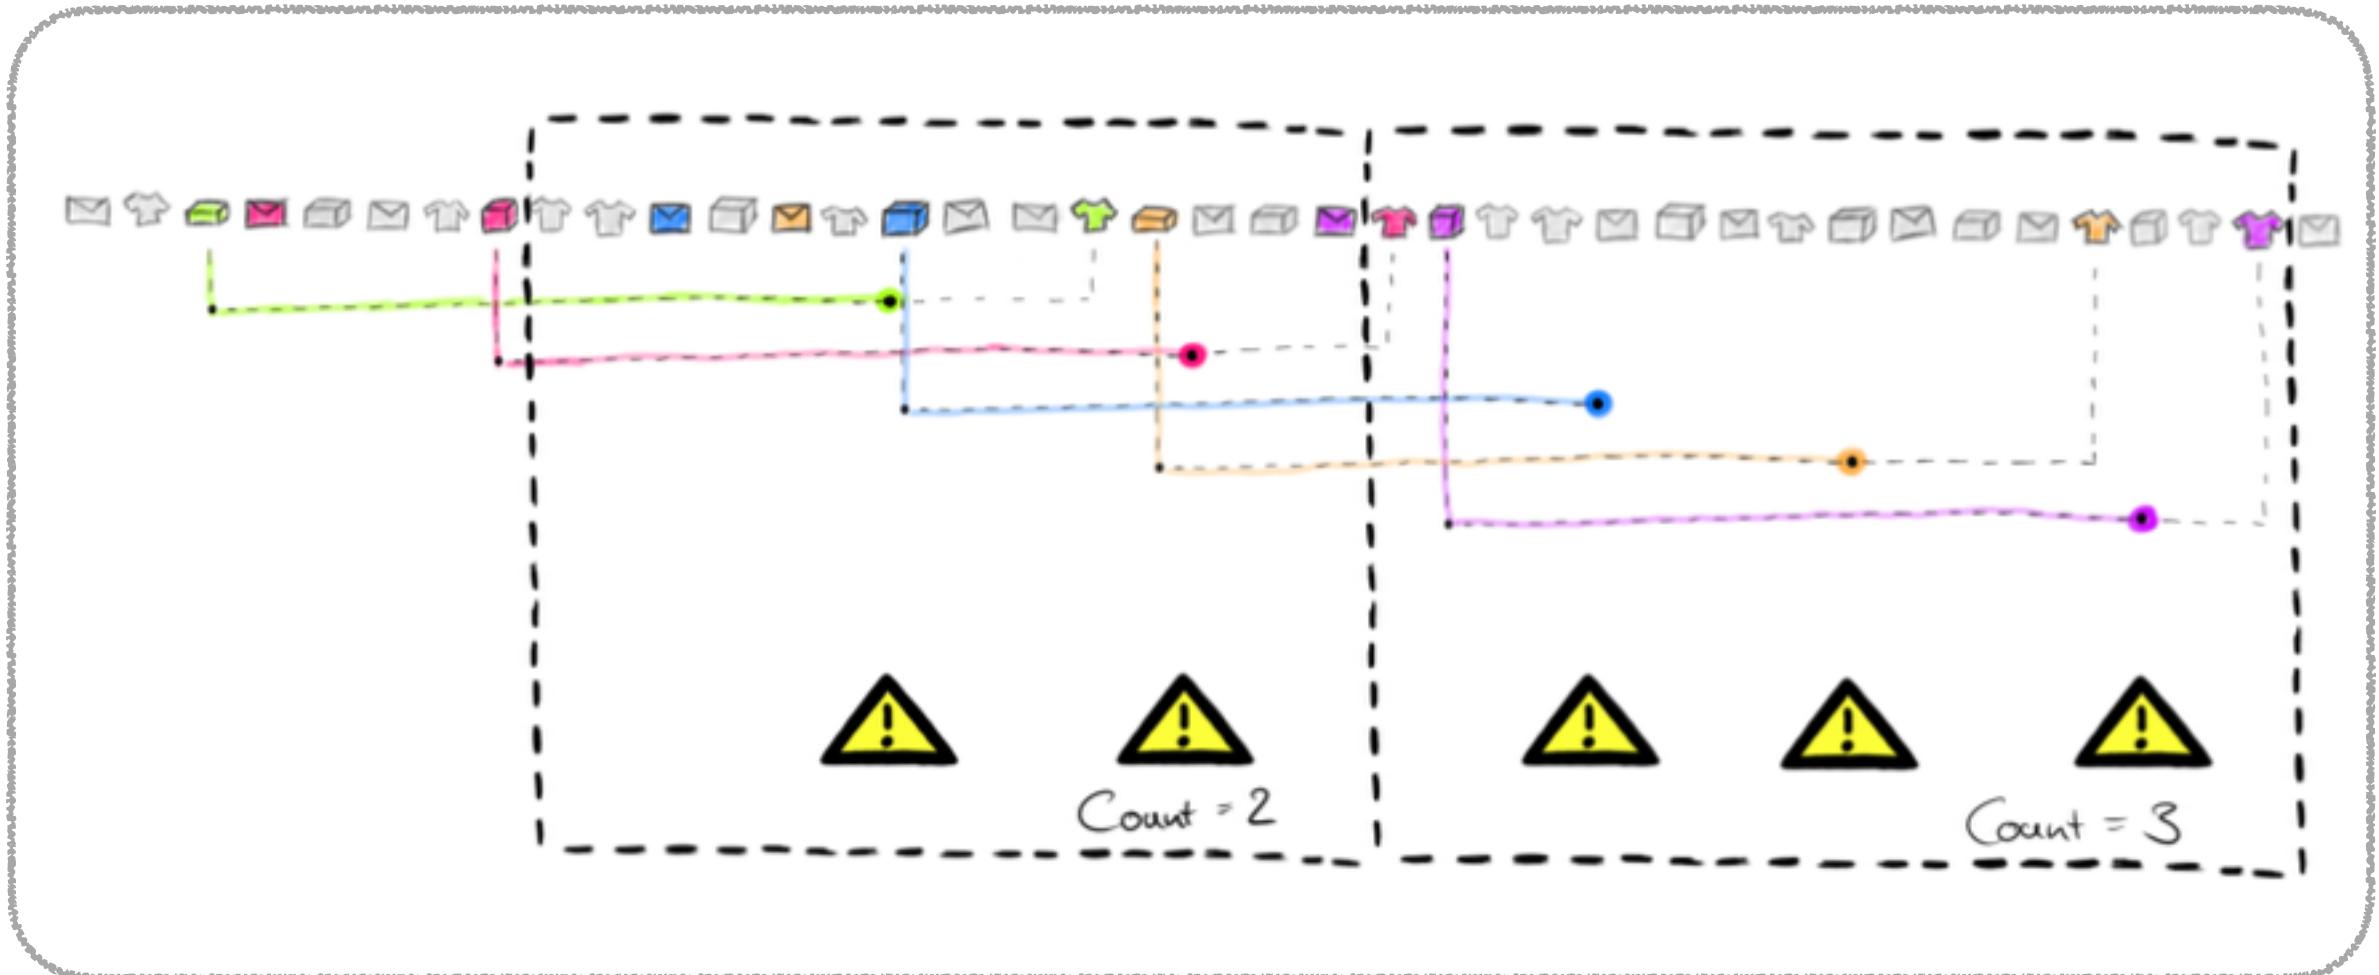
```
val processingPattern = Pattern
  .begin[Event]("received").subtype(classOf[Order])
  .followedBy("shipped").where(_.status == "shipped")
  .within(Time.hours(1))

val processingPatternStream = CEP.pattern(
  input.keyBy("orderId"),
  processingPattern)

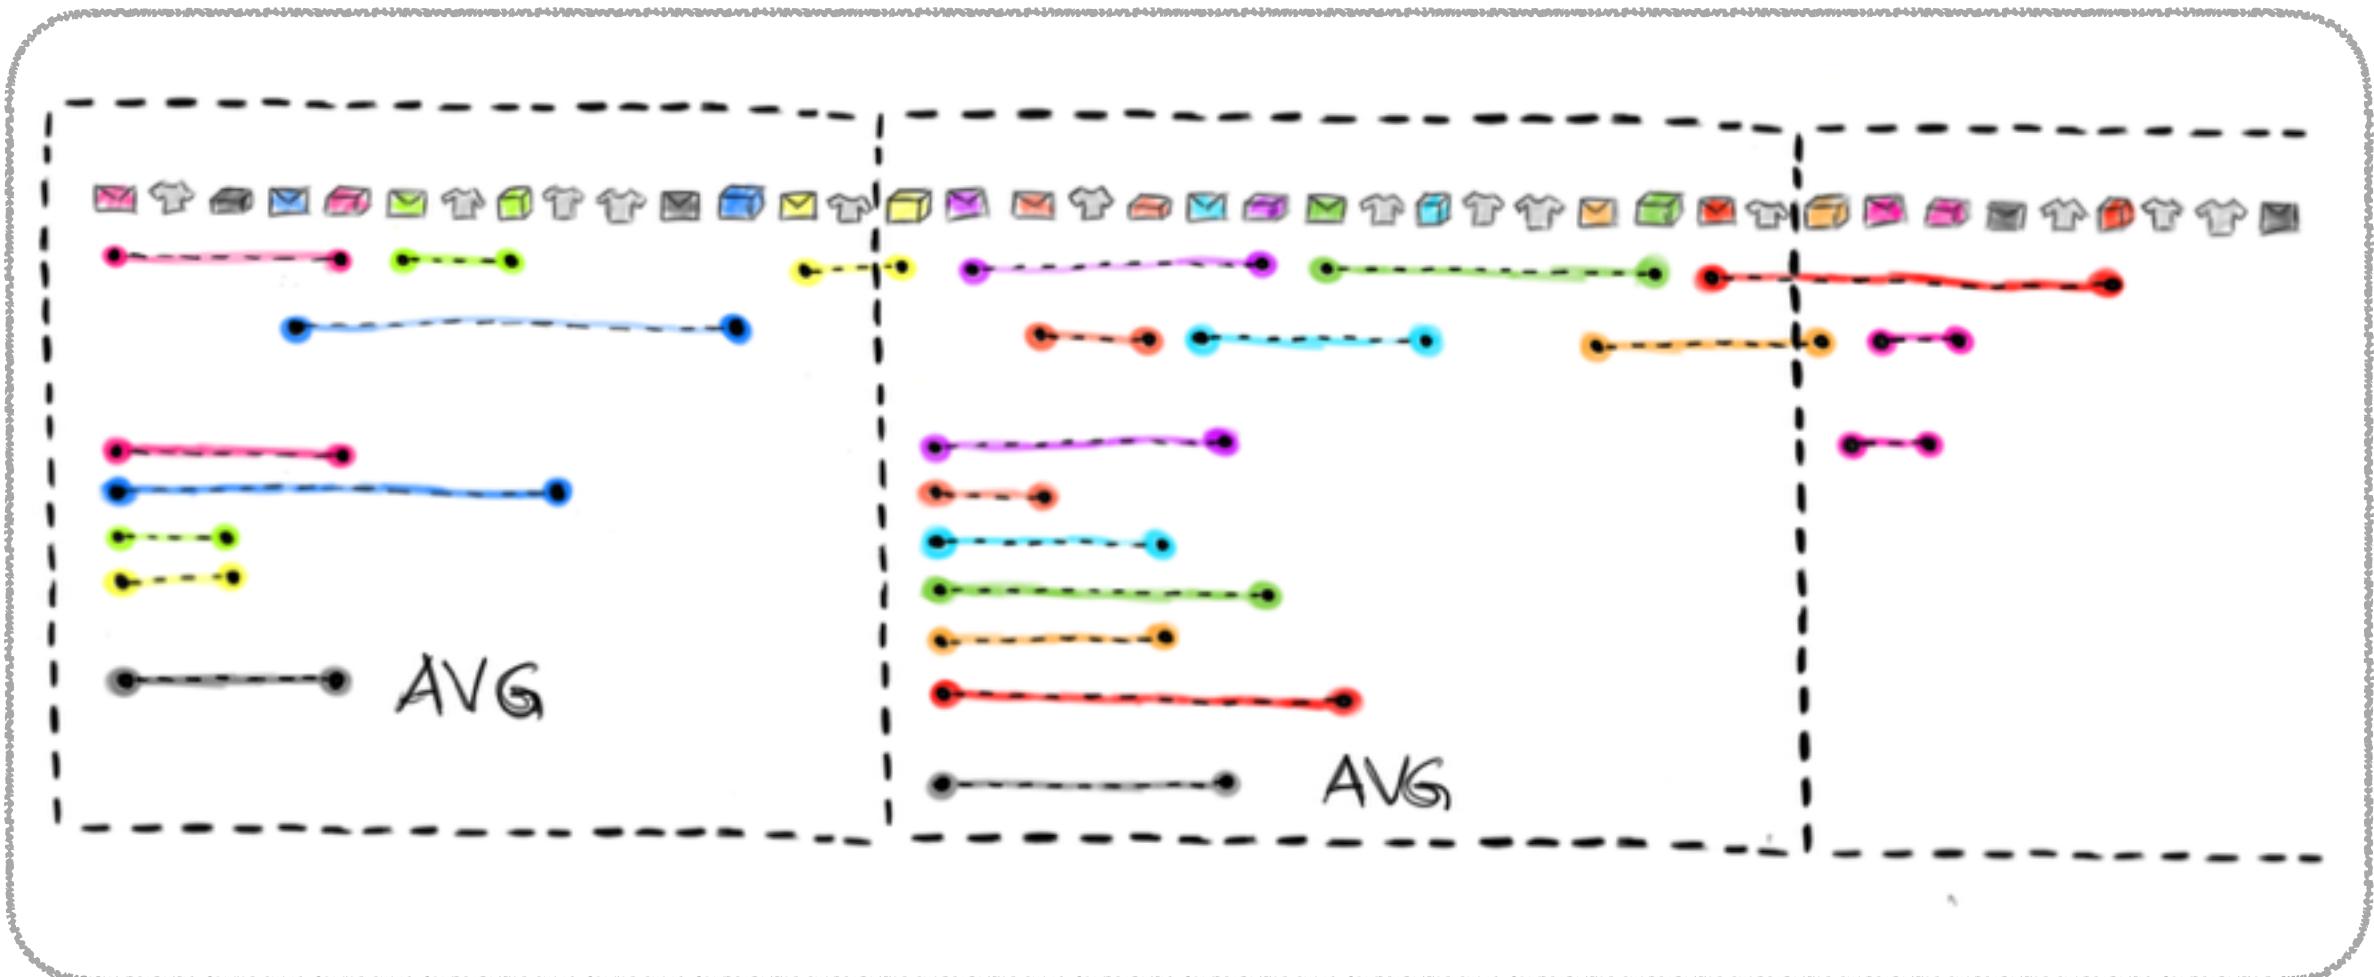
val procResult: DataStream[Either[ProcessWarn, ProcessSucc]] =
  processingPatternStream.select {
    (pP, timestamp) => // Timeout handler
      ProcessWarn(pP("received").orderId, timestamp)
  } {
    fP => // Select function
      ProcessSucc(
        fP("received").orderId, fP("shipped").tStamp,
        fP("shipped").tStamp - fP("received").tStamp)
  }
}
```



Count Delayed Shipments



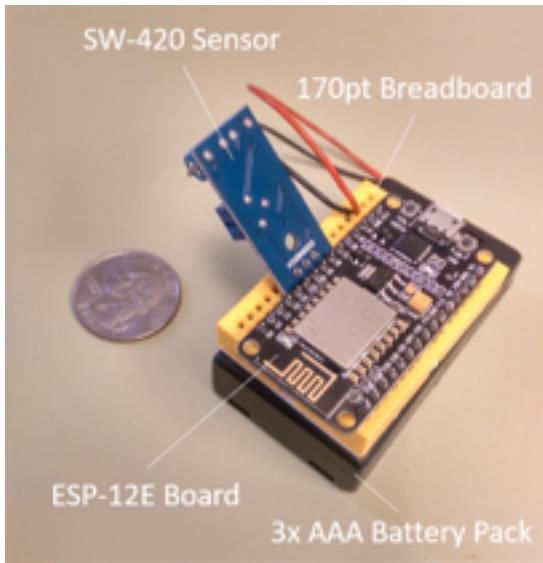
Compute Avg Processing Time



Demonstration

Streaming Analytics

Demonstration



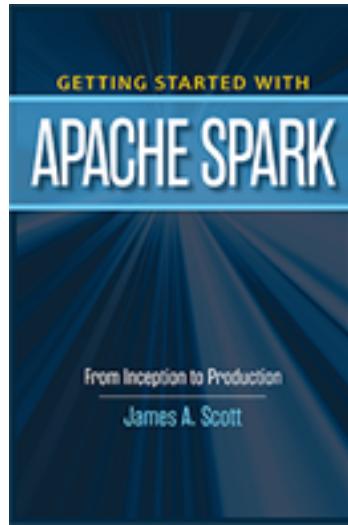
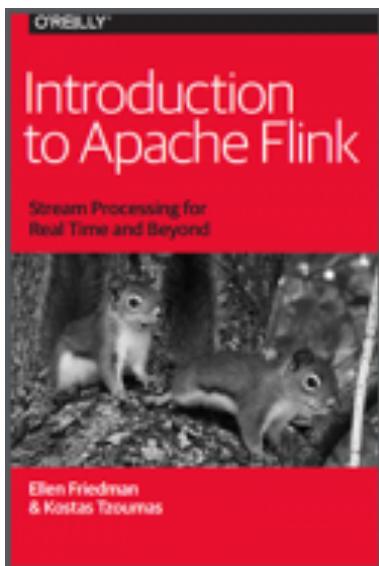
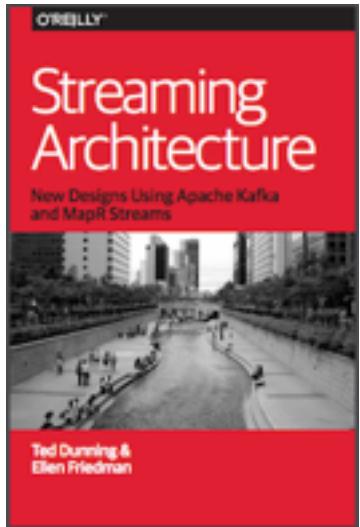
- <https://github.com/mapr-demos/mapr-streams-flink-demo>
- <https://github.com/mapr-demos/wifi-sensor-demo>
- <http://tgrall.github.io/blog/2016/10/12/getting-started-with-apache-flink-and-kafka/>
- <http://tgrall.github.io/blog/2016/10/17/getting-started-with-apache-flink-and-mapr-streams/>
- more soon....

Thanks to



**Kostas Tzoumas
Stephan Ewen
Fabian Hueske
Till Rohrmann
Jamie Grier**

Streaming Architecture



Free ebooks & Online training

<http://mapr.com/ebooks/>

<http://mapr.com/training/>

Stream Processing with Apache Flink

Tugdual Grall
@tgrall

