

# Scaling the tech up

Best Practices, Architecture, Distributed Systems...

Disclaimer

Hi, I'm Julien

# My experience

Streaming, Serverless Computing, PubSub

# What we will cover:

Some basic about distributed systems and scalability

Practice: how we can build a scalable social network

# What we will NOT cover but are equally important:

Scaling the organization

Scaling the tools

Management of the technical debt at Scale

# Distributed Systems

What is it? Is it the cloud (lol)?

Avoid distributing your app as long as you can

Let's define what “scale”  
means in this talk:

The ability to handle more demand!



While maintaining or improving throughput and latency

While maintaining or improving availability and durability. (MeanTimeToRecover should not increase significantly with system size.)

Using infrastructure in a cost-efficient manner: i.e. by adding hardware resources approximately linear in the increase in amount of work

With marginal increase in costs for human operators.

# Scaling up Software

- Is not about using techno, or cloud, or whatever
- It's not free
- It's not magical
- It's distributed systems and data structures and it's engineering excellence, processes and methods.
- **Linus Torvald's** comment: I will, in fact, claim that the difference between a bad programmer and a good one is whether he considers his code or his **data structures** more important. Bad programmers worry about the code. Good programmers worry about **data structures** and their relationships.

# How to scale, the short answer

- Best DS for best operation space/time complexity
- Distribute those DS
- Horizontally scale those distributed DS
- Be careful to:
  - Blast radius
  - Fault tolerance
  - Outages
  - Partitioning (%)

# How to scale, the long answer

- It's a journey:
  - Engineering rigors (code reviews, tests, CM...)
  - CI/CD, Beta, Gamma...
  - Granularity
  - Ops readiness (logs, metrics, alarms) / DevOps / Load Tests
  - Reliable Programming, Data Structures, Define limits
  - PaaS?
  - Distributed Systems (% , Blast Radius)
  - And also: tech debt management, scale up tools, scale up organization...

# Effective method:

- Hope is not a strategy
- Use systemic approach, methods and mechanisms

# Effective methods (example)

1. Go back to the WHY: understanding the business is YOUR business
2. Define precisely what the system should do
  - TLA+!!!
3. Define limits of your systems
4. Use “recovery-oriented computing” at every level
5. Build strong isolation mechanisms to limit blast-radius of problems
6. Prepare for change:

“Definition of Software Architecture: a framework for the disciplined introduction of change” – Tom Demarco

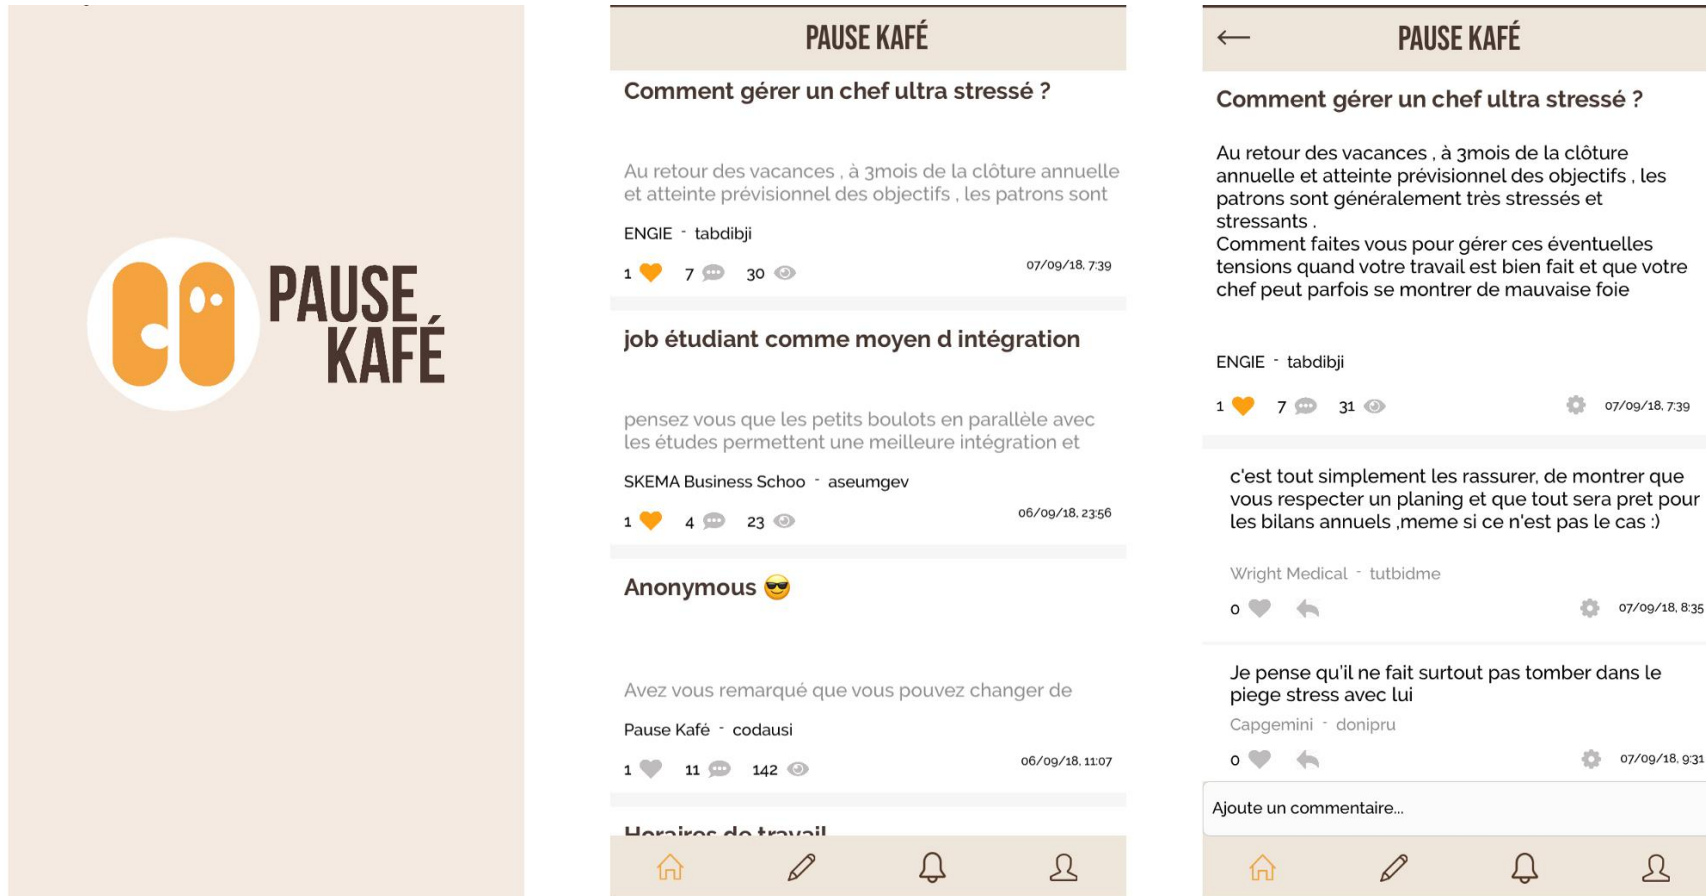
- because: change is inevitable, and without a framework for change, each change results in hacks, and hacks significantly increase the probability of bugs, operator errors, worsening blast-radius, thrash, ...

# Recovery-oriented Computing

- Improve availability by **recovering very quickly** (which you can test), rather than trying to never fail (which is not feasible)
- Aggressively optimize time-to-recover, by:
  - quickly detecting the effects of bugs
  - quickly aborting/killing the smallest “unit of isolation” affected by the bug
  - optimizing the time required to restart each unit of isolation (aka. 'micro-reboots')
- Make your system recover automatically

# Let's try it!

## Let's create a social network: pause-kafe.com





# Data Plane vs Control Plane

# What do we need?

- A feed of posts, like facebook that people can comment. Those comments can also be commented

# But WHY?

- Because it's a new project, new startup, we hope to make a new social network for professional
- We are trying to prove the market

# But WHY?

- Market not proven yet? Take technical debt
- The “Slack” solution; coding isn’t always the answer



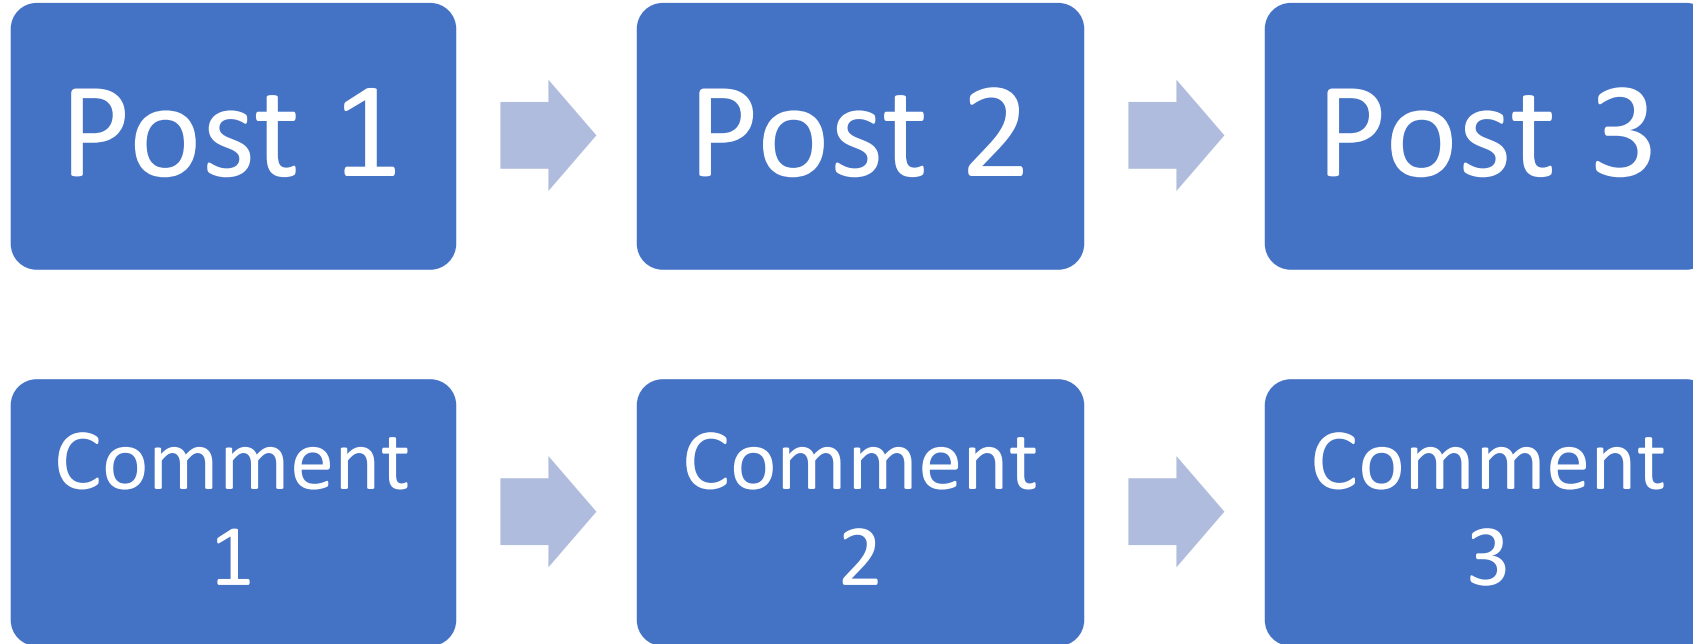
# Define the limits

- Ok now you proved the market and want to do your own dataplane?
- What's the scale you expect?
  - Market size
  - Market share over time (translate into TPS, # of customers...)
  - Size maximum of messages?
  - Review the UX: a huge amount of topic displayed in one shot isn't good
- Size of the team? Not much \$ in bank? Time before company dies? (=low ops?)
- What are the properties of the system (liveness, safety property, non functional):
  - Best ordering
  - Durable
  - Low latency
  - Highly available
  - Can Infinitely grow

# What's the best data structure?

- Operation we will do the most?
  - Get Last Topics
  - Add new Topic
  - Get Next Topics
  - Best time / size complexity? -> Linkedlist!

# What's the best data structure?



# So we need a Linkedlist that can be:

- Low in ops. Most likely need a PaaS
- Distributed
- Horizontally scalable
- Replicated (fault tolerance, availability and co)
- Stored on disk
- Dead-lock free
- Space/cost optimized



# Hashtable: the mother of all

- With an hashtable you can implement all DS in the world (in mostly the good ways)
- A good distributed hashtable – cloud scale ready will solve most the technical challenges
- Basic operations we need:
  - CAS !!!
  - Atomic append
  - Batch write and read
  - Atomic updates with conditional writes
  - Increments without read
  - Etc...

# Hashtable: the mother of all

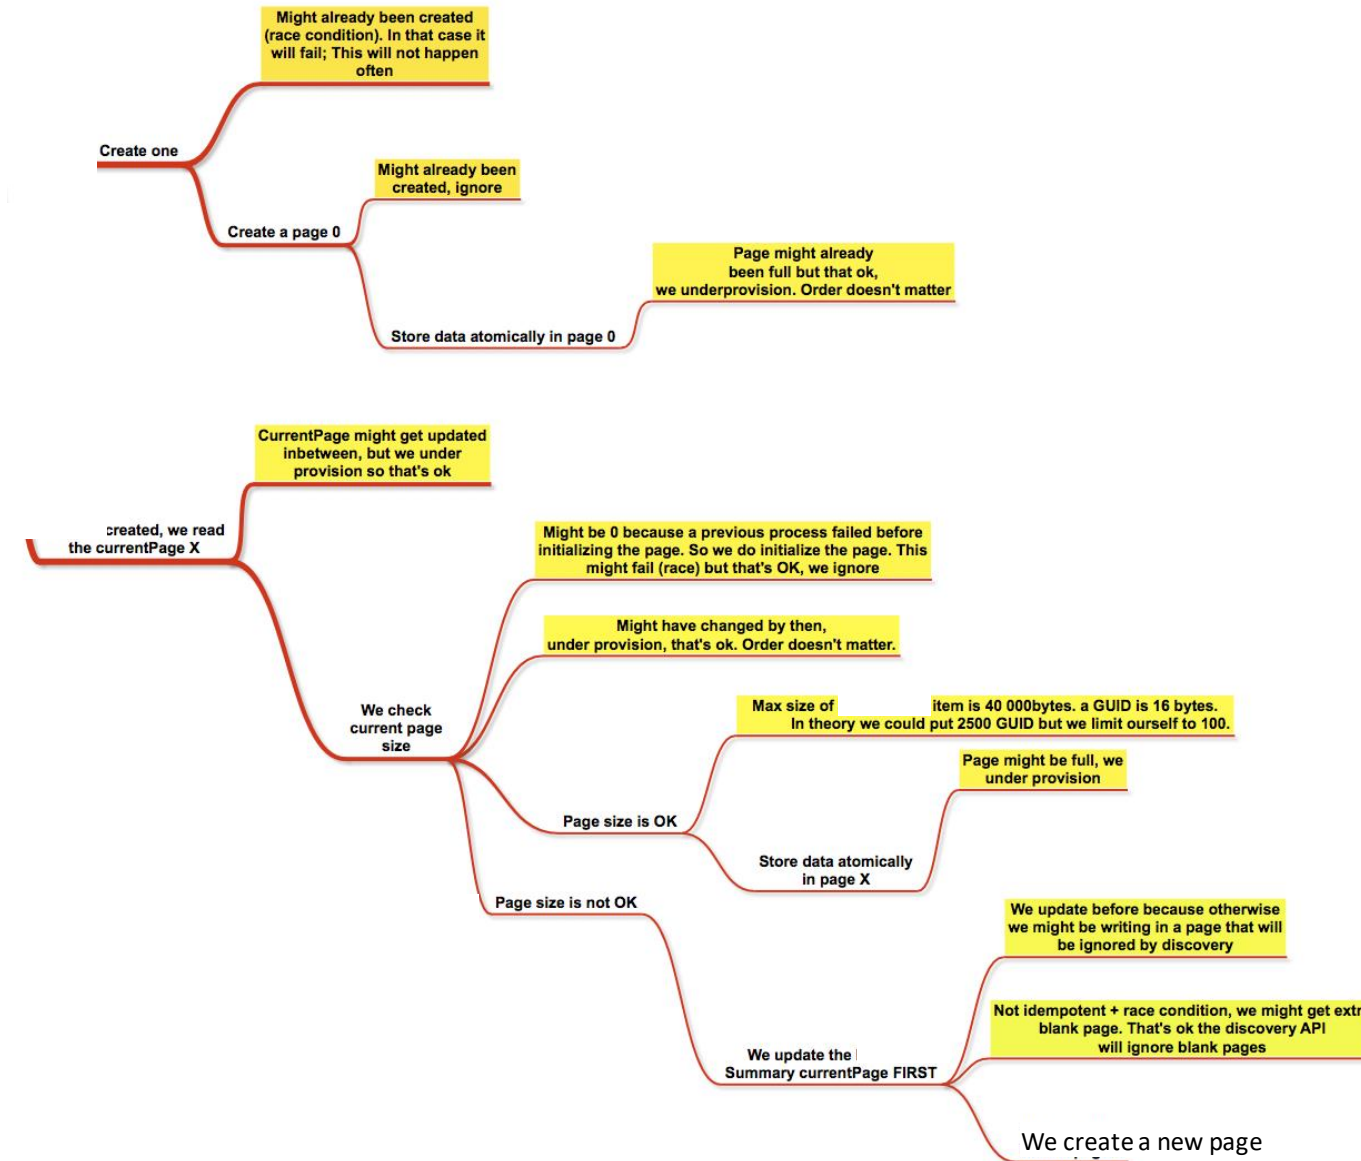
- OCI Object Store MD, DynamoDB, CosmosDB etc... allows to store key value efficiently with strong consistency and consistent partitioning on the primary key.
- We have the need to represents the data in form of a linked-list, and we want to benefits from the same guarantee as those KeyValueDB (availability, consistency, latency...).
- So let's build a facade on top that allow us to have our Linkedlist
- A linked list is represented by multiple "page".
- The first page of the LinkedList will contain a pointer to the current last "page" of the linkedlist (where new items should be appended), it is called the "summary" page.
- The next pages, also called "data page" contains the data appended to the linkedlist. Each page can contain an array of item and can store between 0 to N item.
- The goal of putting together those items in the same node is to not have to create a new element for each writes. Why? because we want to avoid to have to update the pointer in the summary page too often. This avoid doing too much random access write on the summary page, which on top of being slower, will be a "hotspot" prone to bug at scale. N threads using the linkedlist could increment the page summary in the same time which would lead to a lot of "blank" page.

# Our distributed Linked List



# Recovery programming

- What wrong could happen when append data for example ? (bug, race conditions etc)



Some code

/\* This function is to optimize how much a data can be added by “page”.

If page == 1 (which is the minimum) it will create a new keyvaluestoredb item for each node in the linkedlist.

This will most likely not react well at scale (empty page, pointer incremented too much...).

The goal is to put as much data as possible in a single page but:

- The size should not go being the maximum keyvaluestore item size (which is 400 KB)

- BulkRetrieve in keyvaluestore is limited to 16 000KB. You will most likely want to retrieve a linkedlist in bulk, so keep this in mind.

- Each page contains 1KB of metadata

- Concurrent threads can be trying to add in the same page despite the limit being reached! It should be small window of few seconds when this happen, but my advice would be to over provisioning by assuming the maximum size of the keyvaluestore DB item is divided by 2 (and then tune things up or down depending on the scale)

- Overall, the number you put here, is the maximum of TPS for this instance of the linkedlist (req in the average same seconds) you can have before any of your customers get TooBigItemException

So for example, if the maximum size of value is 1 000 characters, in UTF-8, it mean each element maximum 1,074 bytes (1kb).

The maximum size of a page should be:  $1\text{KB} + 1\text{kb} * N \leq 400\text{KB}$ , N should be less than 390ish.

Then we apply the over-provisioning factor of /2, so we limit ourself to 200 elements.

Note: it would limit to retrieve 80 linked-list page in a BulkRetrieve operation \*/

```
module.exports.configureMaximumNumberOfElementPerPage = function(numberOfMaximumElement) {  
    if (numberOfMaximumElement) {  
        config.maxElementPerPage = numberOfMaximumElement;  
    }  
}
```

```
/* This will create the first summary page. This operation is idempotent. */  
module.exports.idempotentCreate = function(id, metadata) {  
  const summaryObject = defaultPageSummary(id, metadata);  
  
  const itemInfo = {  
    Item: summaryObject,  
    Condition: '!summaryObject.id', //CAS on the id for idempotency  
  };  
  
  return keyValueStoreDb.put(itemInfo).promise().then(res => summaryObject);  
}
```



```
/*
```

If the page does not exist it first atomically increments the pointer to point to the “current” page.

This pointer might be incremented at the same time as another thread which is why you can have “blank” pages that you should ignore.

It then creates the new page by using the CAS mechanism. If the page has been created in between, it just fails silently.

Then it finally appends the value (or directly appends if the page was already created) by inserting the data into the page atomically.

This operation is not idempotent and could insert duplicates due to retries.

Note: value needs to be an object

```
*/
```

```
module.exports.atomicAppend = function(id, value) {  
  return getCurrentPage(id).then(function(currentPage) {  
    return atomicAppendImpl(id, currentPage, value);  
  });  
}
```

/\* It will retrieve the top N most recent item that has been appended to the linkedlist.

Be careful to bufferoverflow here. Avoid asking for a crazy amount \*/

```
module.exports.retrieveLastMostRecent = function(id, numberOfItems, callback) {  
    currentPage(id)  
    .then(function(currentPage) {  
        retrieveNElement(id, currentPage, null, numberOfItems, callback);  
    });  
}
```

/\* It will retrieve the next N most recent item that has been appended to the linkedlist.

To use this, you need to specify where the algorithm should start reading.

Each object retrieved from the Linkedlist contains a PageId, and a SequenceId, this is what you need to use to generate the pointer to dictate where to start reading.

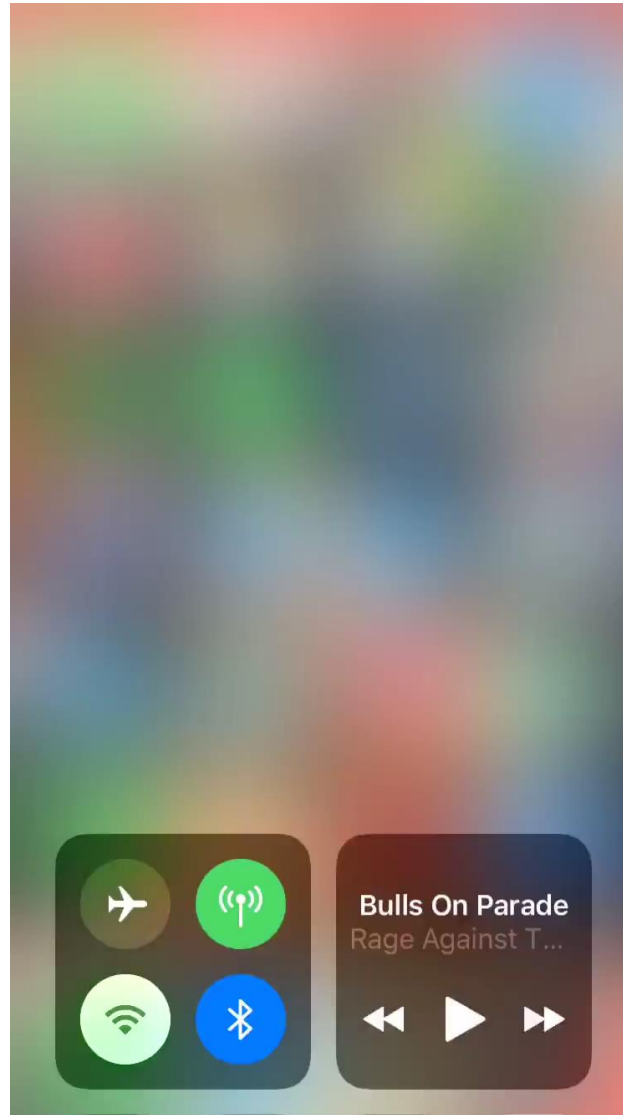
\*/

```
module.exports.retrieveNextMostRecent = function(id, startAfterPointer, numberOfItems, callback) {  
    if (!startAfterPointer ||  
        !startAfterPointer.page_id ||  
        !startAfterPointer.sequence_id) {  
        throw 'No valid pointer has been set';  
    }  
  
    const pointer = getValidPointer(startAfterPointer);  
    retrieveNElement(id, pointer.page_id, pointer.sequence_id, numberOfItems, callback);  
}
```

# So what do we have?

- A Linkedlist that is distributed, horizontally scalable, best ordering, dead-lock free, space/cost optimized, Linked List built on top of a key value store
- We know the limits of it (and it is configurable)
- We can most likely build Facebook, Twitter, any feed news on top of it.
  - I've load tested; 100 concurrent users in the same 25ms max per Linkedlist. Horizontally scale to billions of linkedlists (consistent hashing FTW)
- It would be contained in a single dependency (easy to maintain)
- Only 427 lines of code
- Most the hard things is not our problem but our PaaS provider problem: WIN
- And we never talked about technology 😊

# Video



Last advices

Let the complex distributed  
systems thingy be owned by  
someone else trustworthy  
(hardware, replication,  
consensus etc...)

Stay Simple  
Be a hacker  
But do everything well

Use Science  
(data structures, TLA+ and  
co)



Use recovery  
programming and  
Operations readiness day 1

Have a plan, understand  
that scale is a journey

Be rigorous

Q&A